



L7: Memory Basics and Timing



Acknowledgement: Nathan Ickes, Rex Min

**J. Rabaey, A. Chandrakasan, B. Nikolic, “Digital Integrated Circuits: A Design Perspective”
Prentice Hall, 2003 (Chapter 10)**



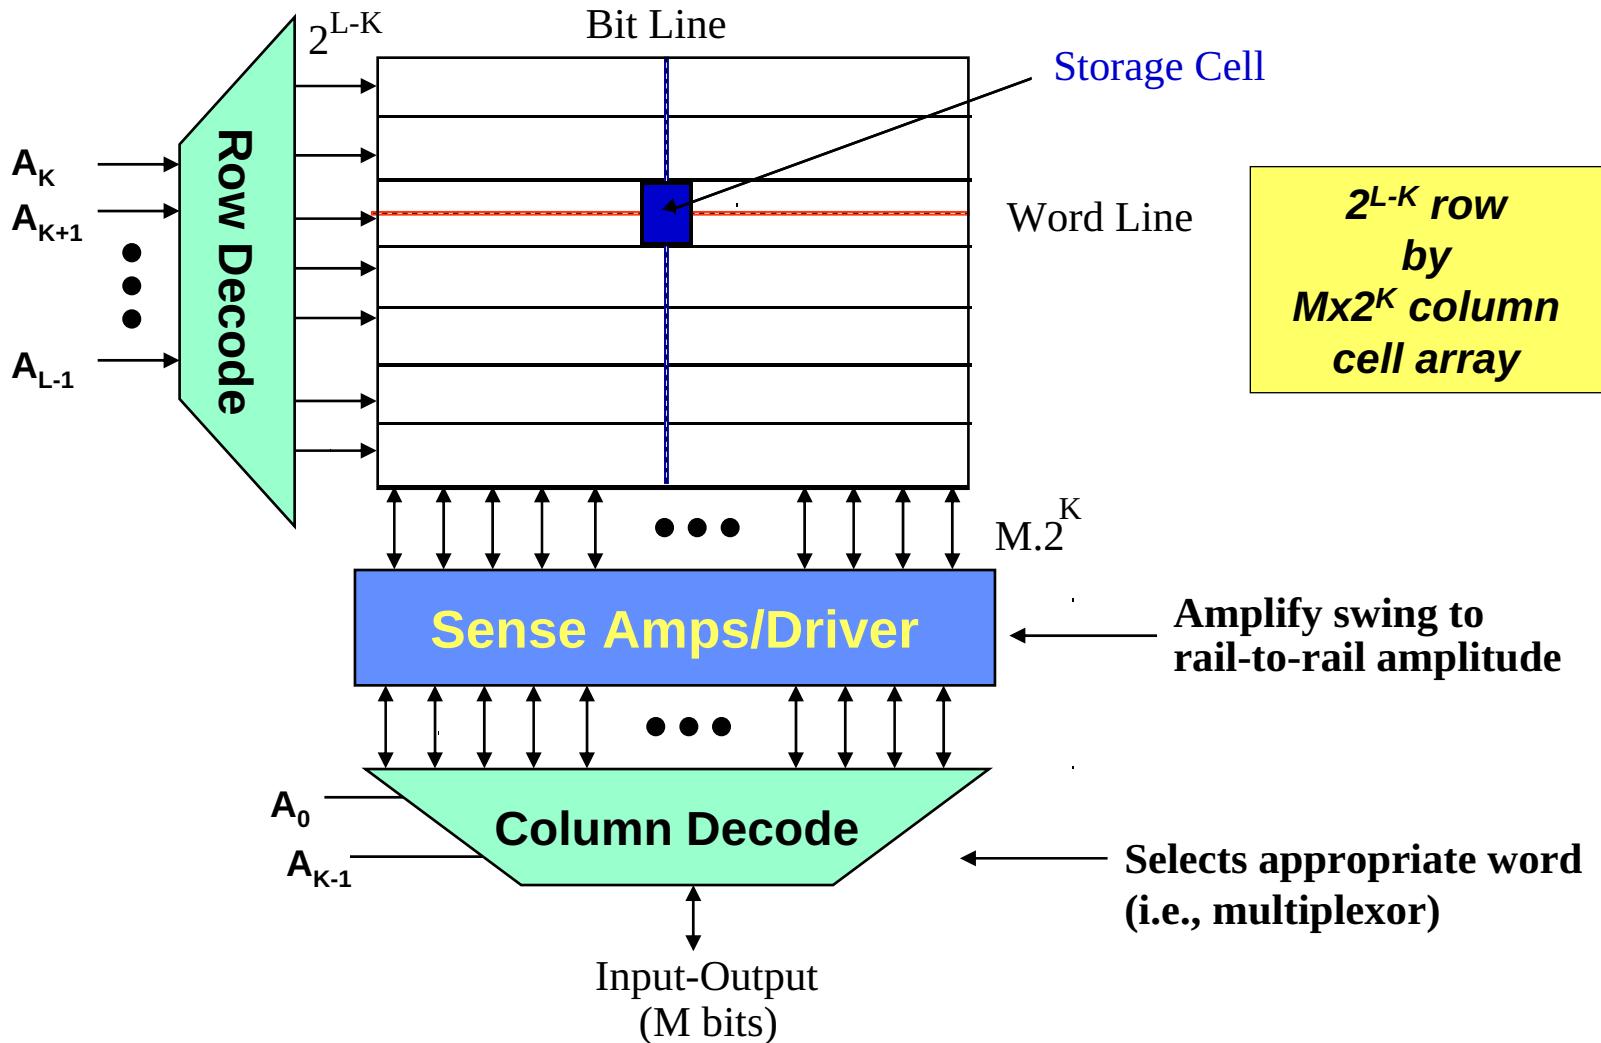
Memory Classification & Metrics

Read-Write Memory		Non-Volatile Read-Write Memory	Read-Only Memory (ROM)
Random Access	Non-Random Access	EPROM E ² PROM FLASH	Mask-Programmed
SRAM	FIFO LIFO		

Key Design Metrics:

1. Memory Density (number of bits/ μm^2) and Size
2. Access Time (time to read or write) and Throughput
3. Power Dissipation

Memory Array Architecture

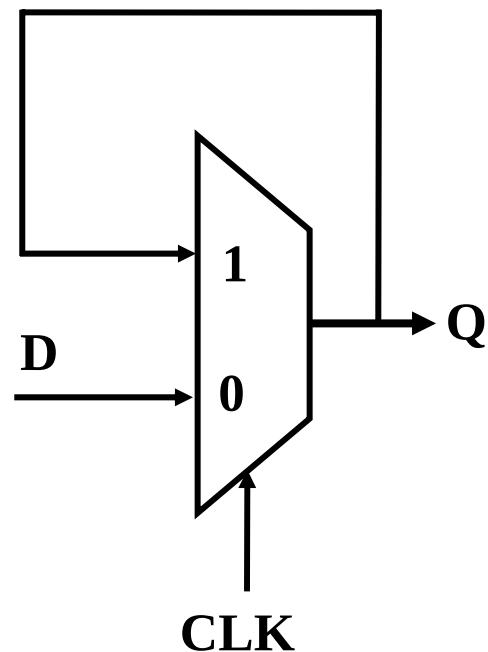
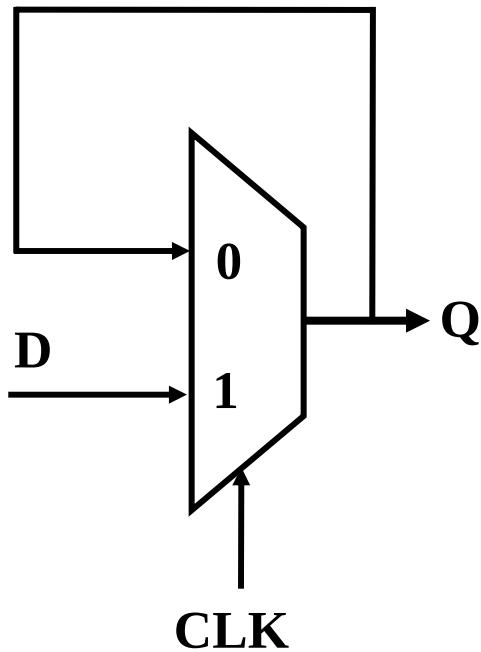




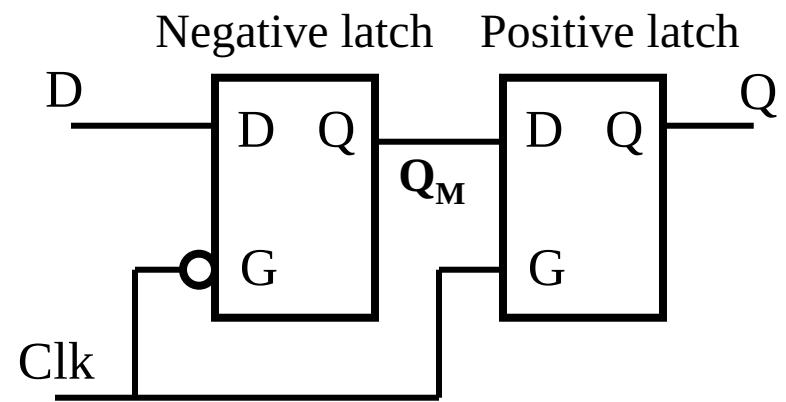
Latch and Register Based Memory



Positive Latch Negative Latch



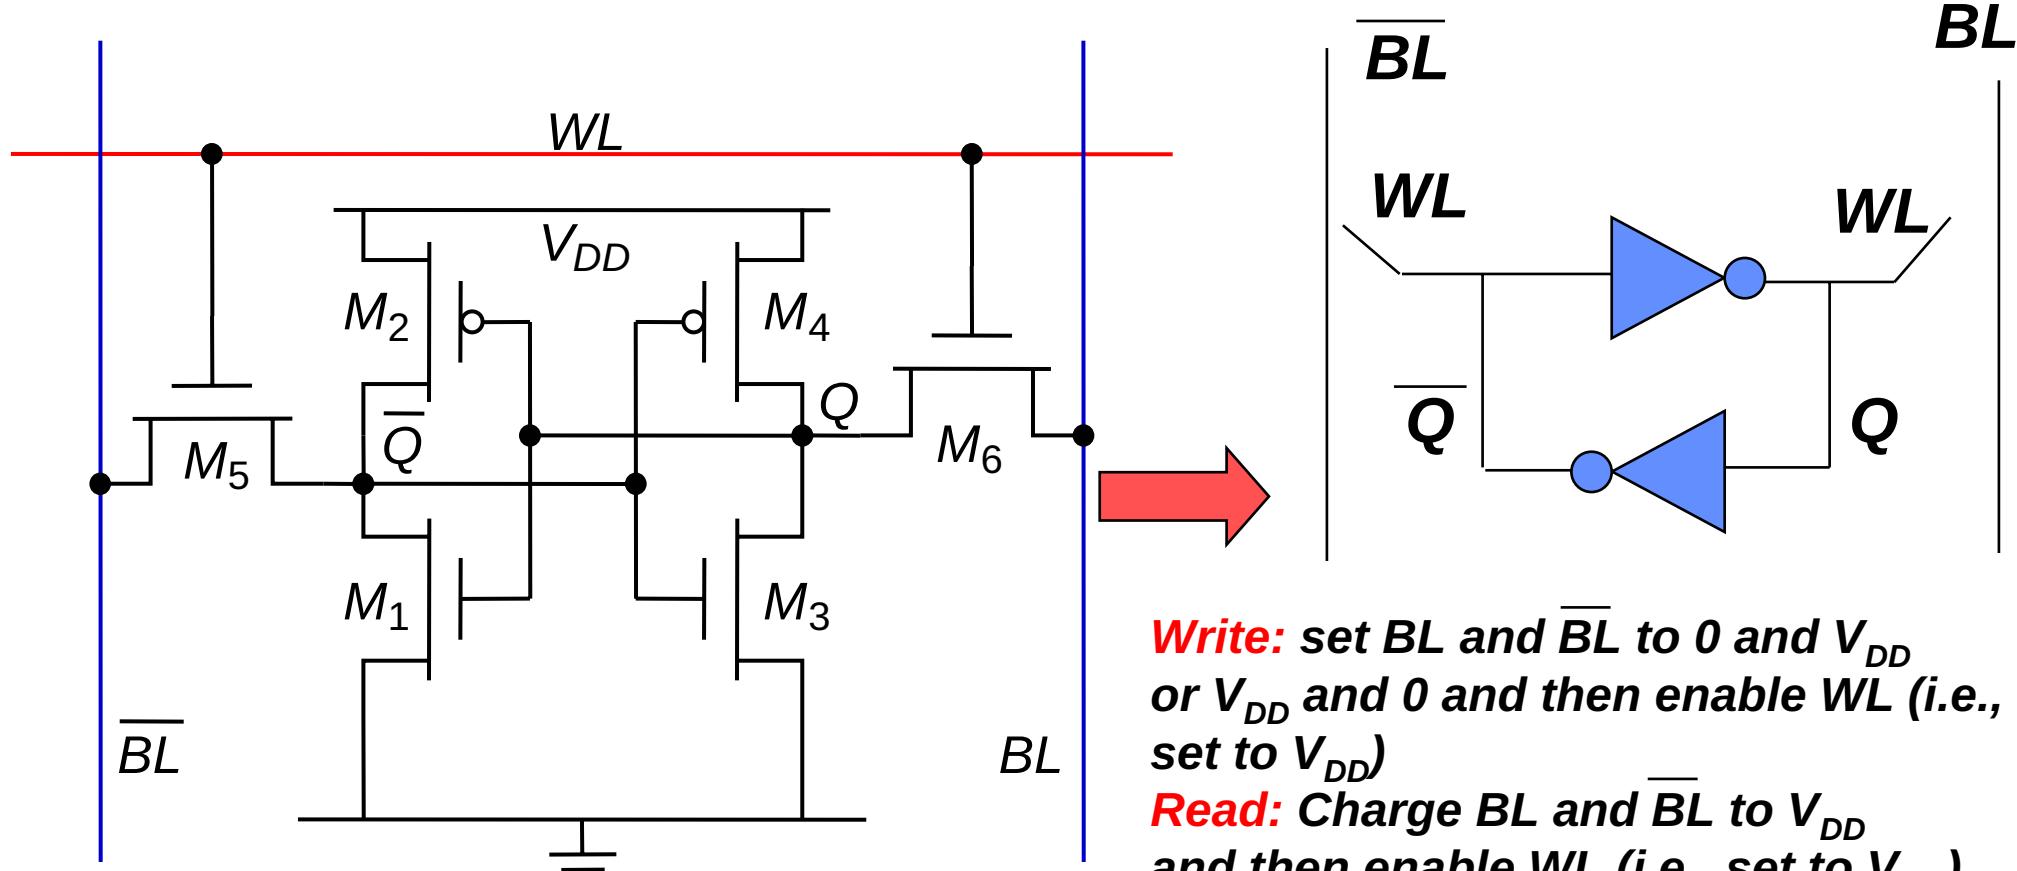
Register Memory



- **Works fine for small memory blocks (e.g., small register files)**
- **Inefficient in area for large memories – density is the key metric in large memory circuits**

How do we minimize cell size?

Static RAM (SRAM) Cell (The 6-T Cell)

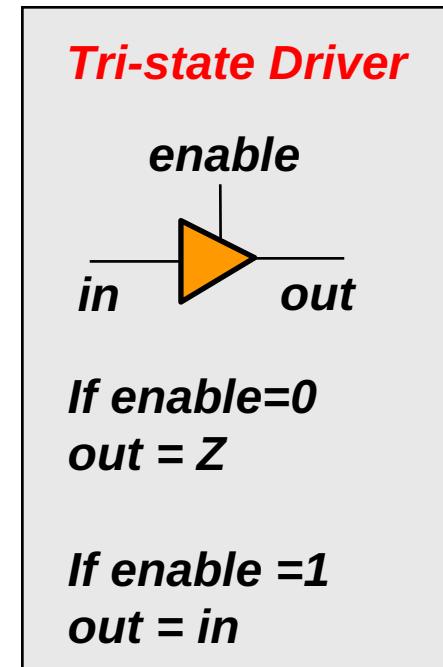
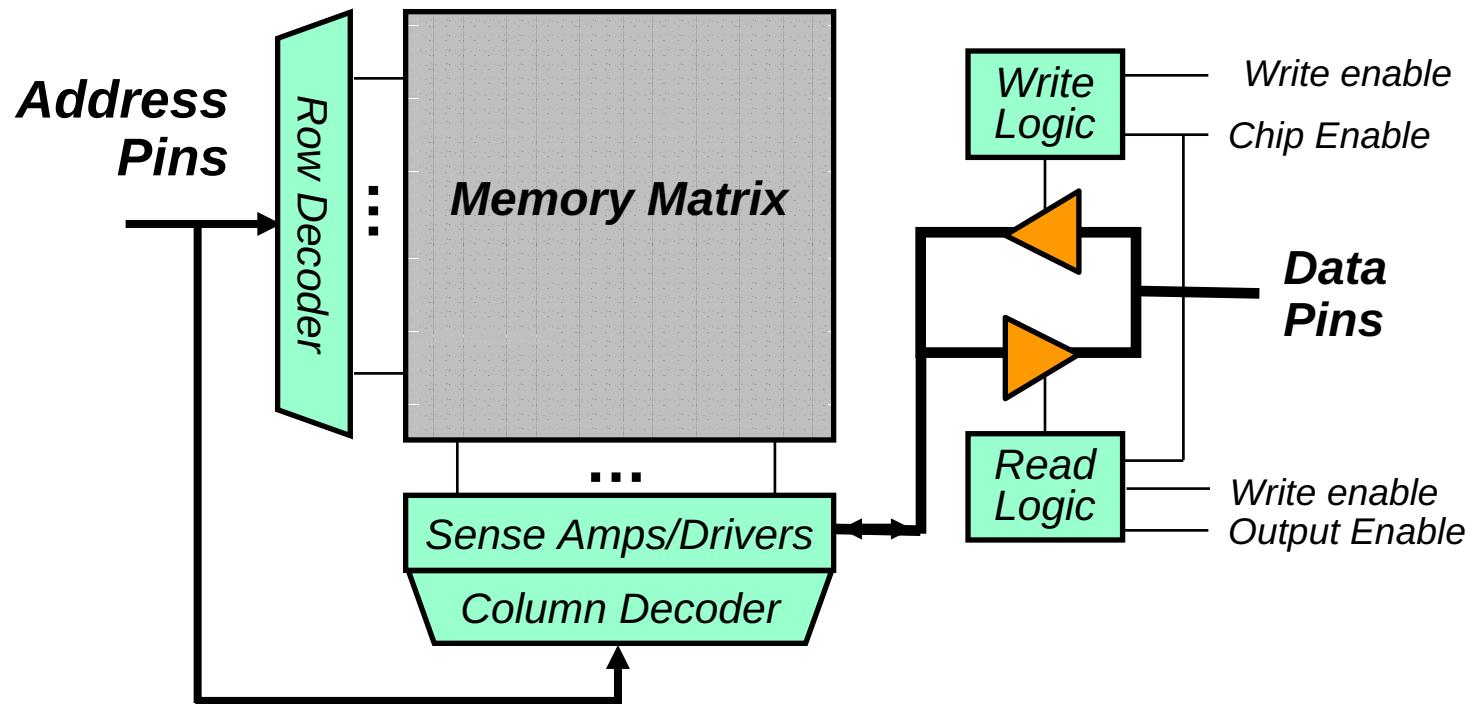


Write: set BL and \bar{BL} to 0 and V_{DD} or V_{DD} and 0 and then enable WL (i.e., set to V_{DD})

Read: Charge BL and \bar{BL} to V_{DD} and then enable WL (i.e., set to V_{DD}). Sense a small change in BL or \bar{BL}

- State held by cross-coupled inverters (M₁-M₄)
- Retains state as long as power supply turned on
- Feedback must be overdriven to write into the memory

Interacting with a Memory Device



- **Address pins** drive row and column decoders
- **Data pins** are bidirectional and shared by reads and writes

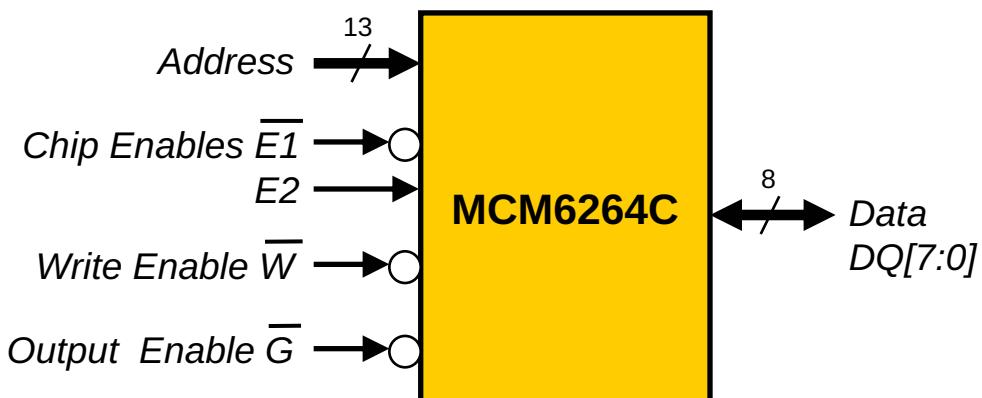
- **Output Enable** gates the chip's tristate driver
- **Write Enable** sets the memory's read/write mode
- **Chip Enable/Chip Select** acts as a “master switch”



MCM6264C 8k x 8 Static RAM

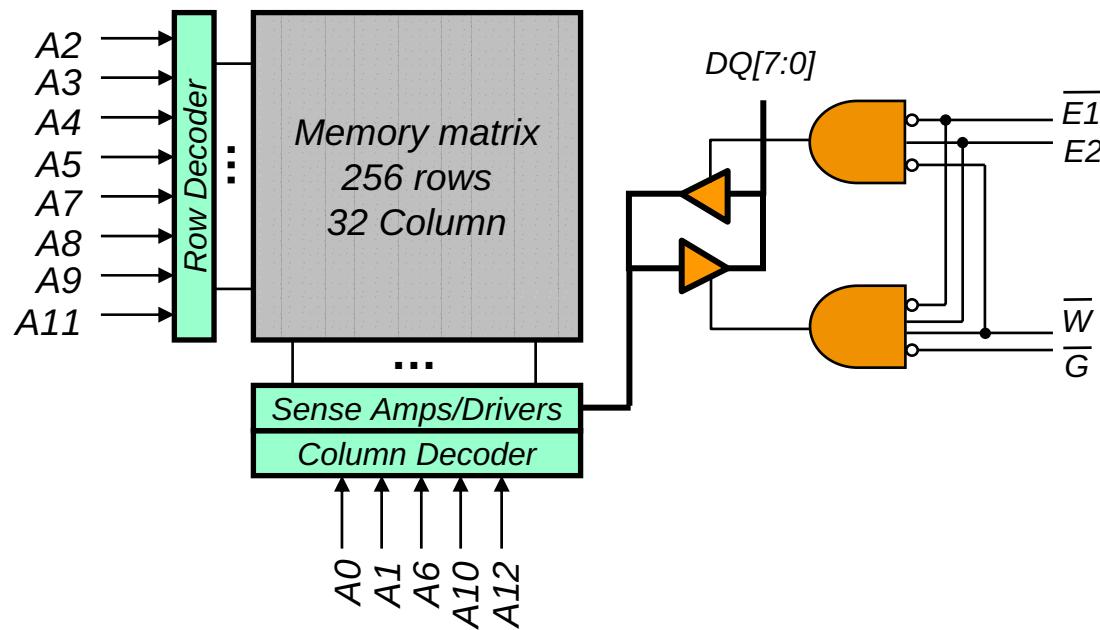


On the outside:



- Same (bidirectional) data bus used for reading and writing
- Chip Enables ($\bar{E}1$ and $E2$)
 - $E1$ must be low and $E2$ must be high to enable the chip
- Write Enable (\bar{W})
 - When low (and chip is enabled), the values on the data bus are written to the location selected by the address bus
- Output Enable (\bar{G})
 - When low (and chip is enabled), the data bus is driven with the value of the selected memory location

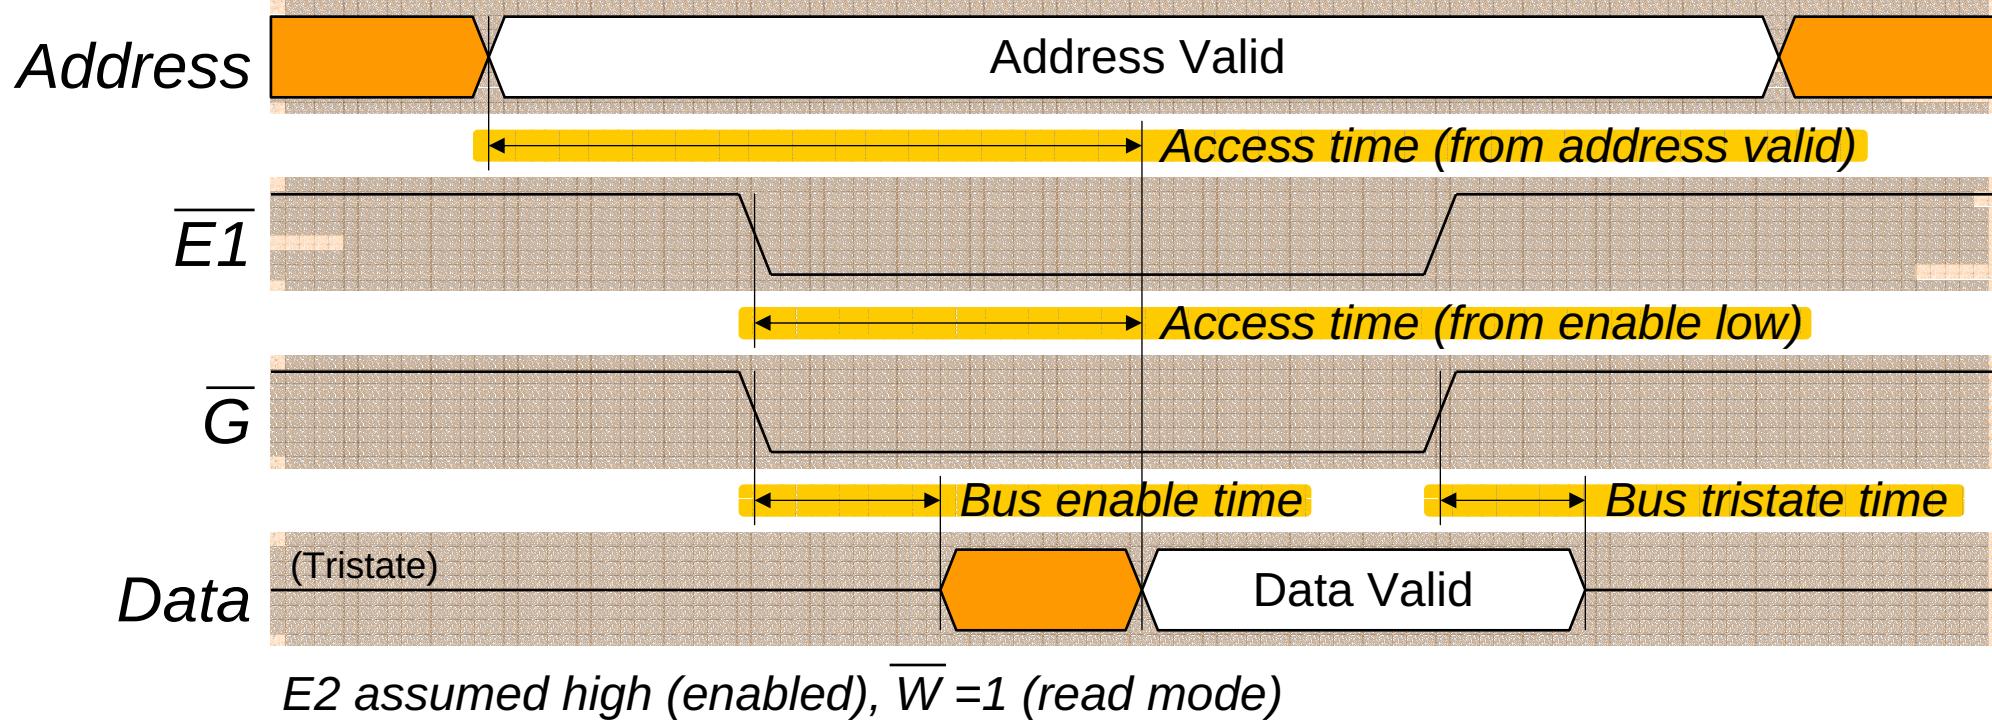
On the inside:



Pinout

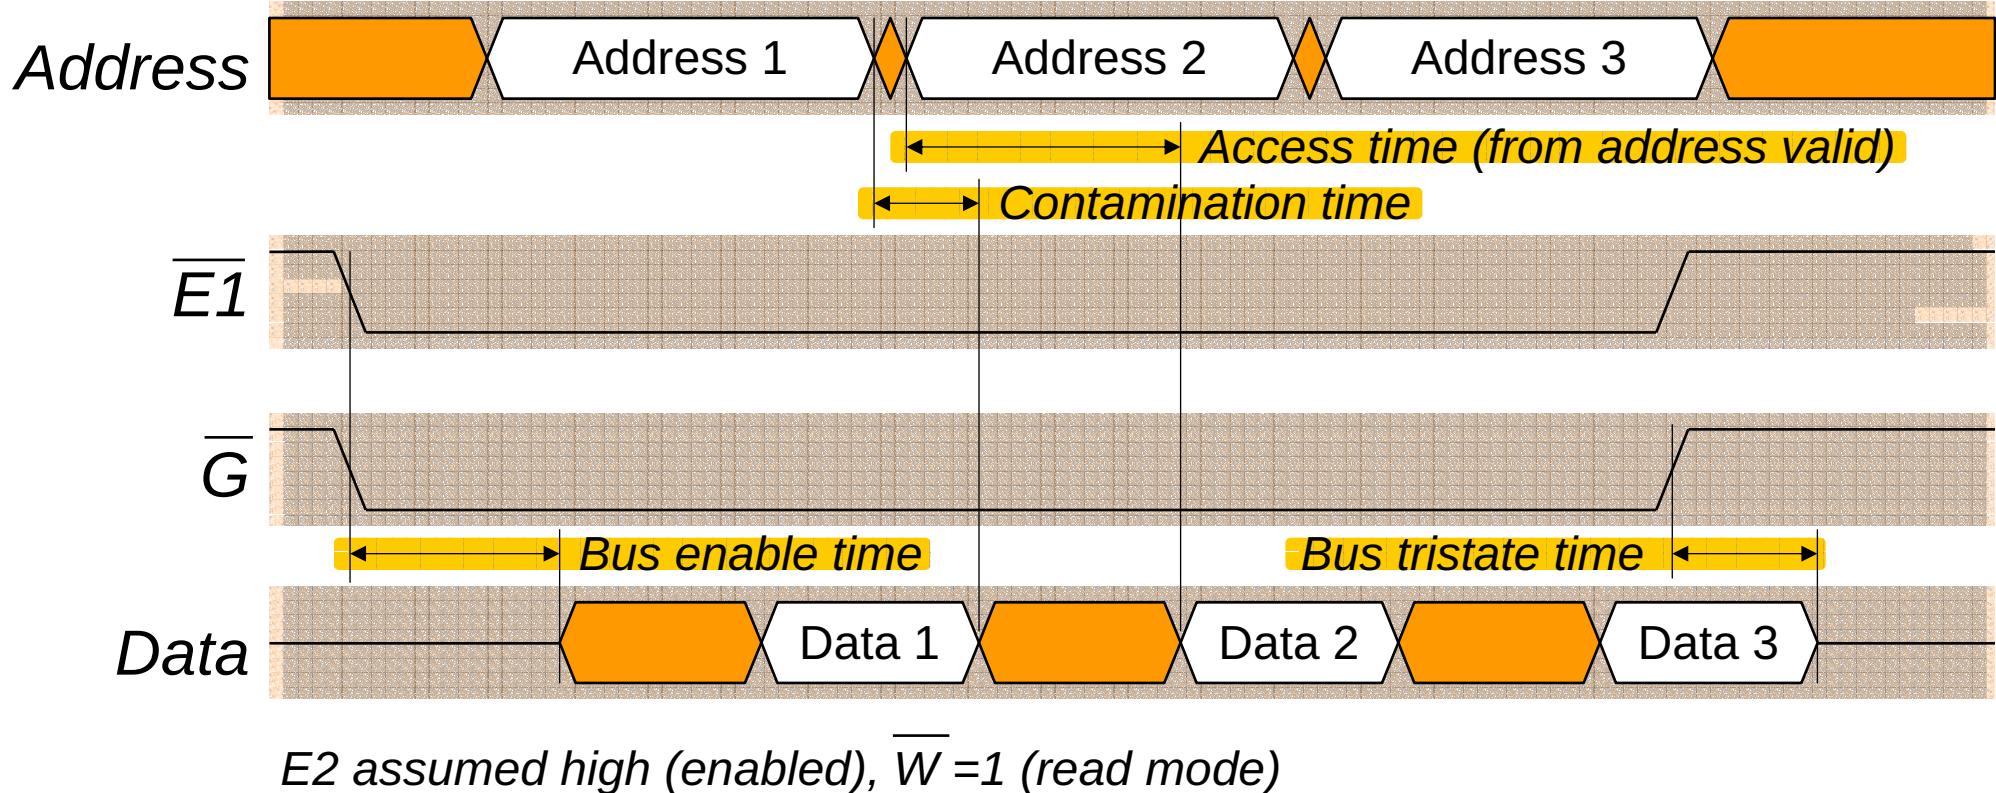
NC	1	V _{CC}
A12	2	\bar{W}
A7	3	$E2$
A6	4	A8
A5	5	A9
A4	6	A11
A3	7	\bar{G}
A2	8	A10
A1	9	$E1$
A0	10	DQ7
DQ0	11	DQ6
DQ1	12	DQ5
DQ2	13	DQ4
V _{SS}	14	DQ3

Reading an Asynchronous SRAM



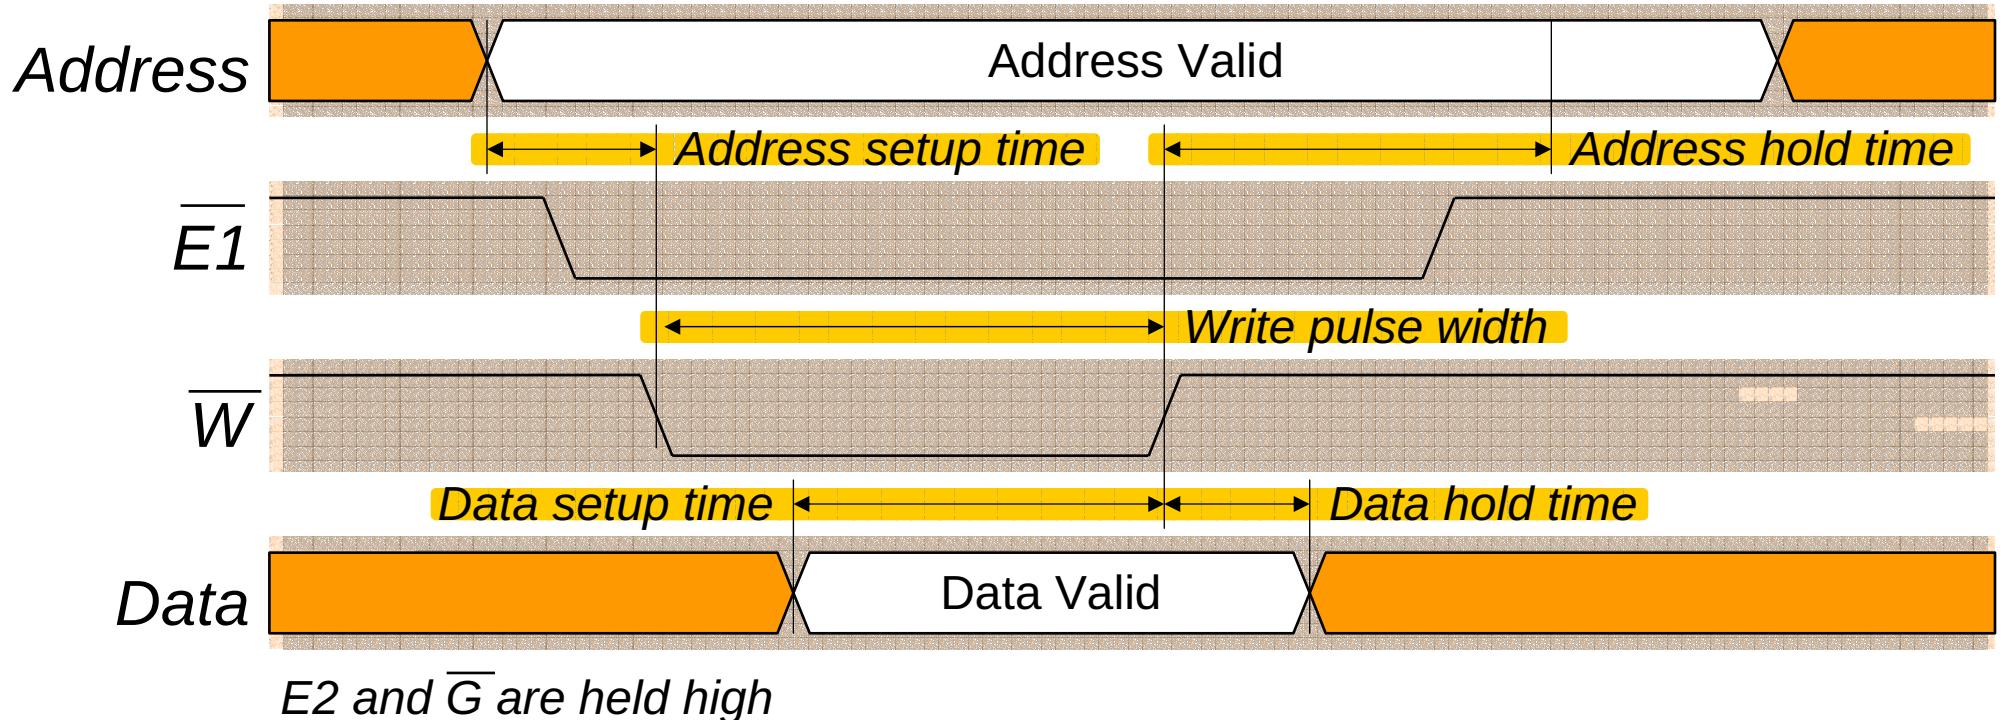
- Read cycle begins when all enable signals ($\overline{E1}$, $\overline{E2}$, \overline{G}) are active
- Data is valid after read access time
 - Access time is indicated by full part number: MCM6264CP-12 → 12ns
- Data bus is tristated shortly after \overline{G} or $\overline{E1}$ goes high

Address Controlled Reads



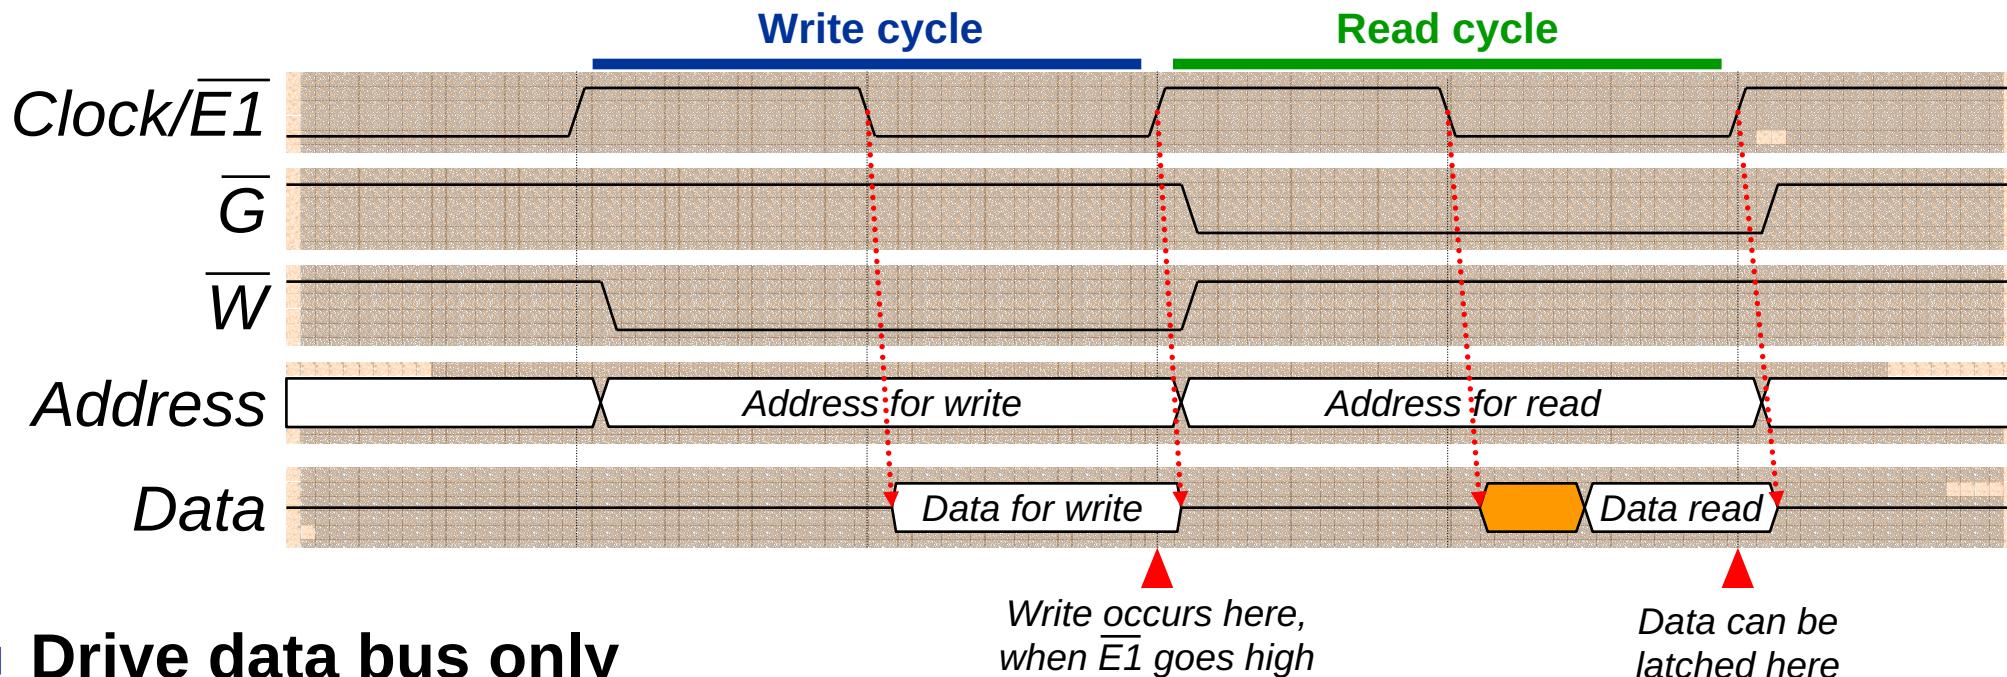
- Can perform multiple reads without disabling chip
- Data bus follows address bus, after some delay

Writing to Asynchronous SRAM

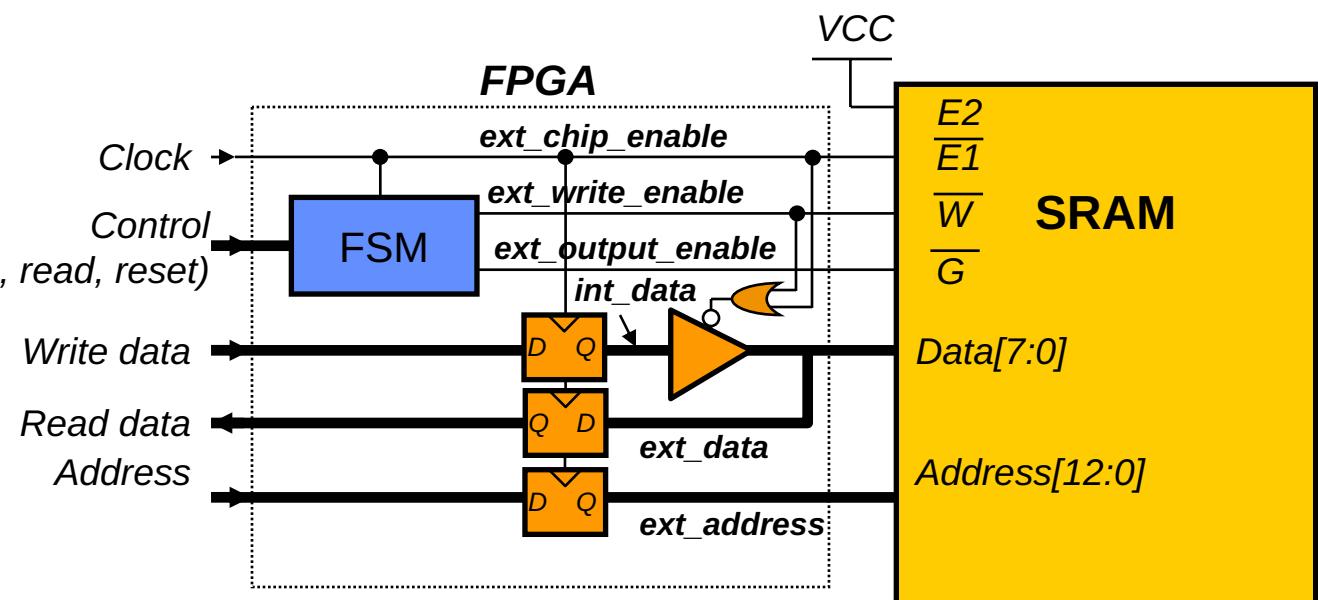


- **Data latched when \overline{W} or $\overline{E1}$ goes high (or $E2$ goes low)**
 - Data must be stable at this time
 - Address must be stable before \overline{W} goes low
- **Write waveforms are more important than read waveforms**
 - Glitches to address can cause writes to random addresses!

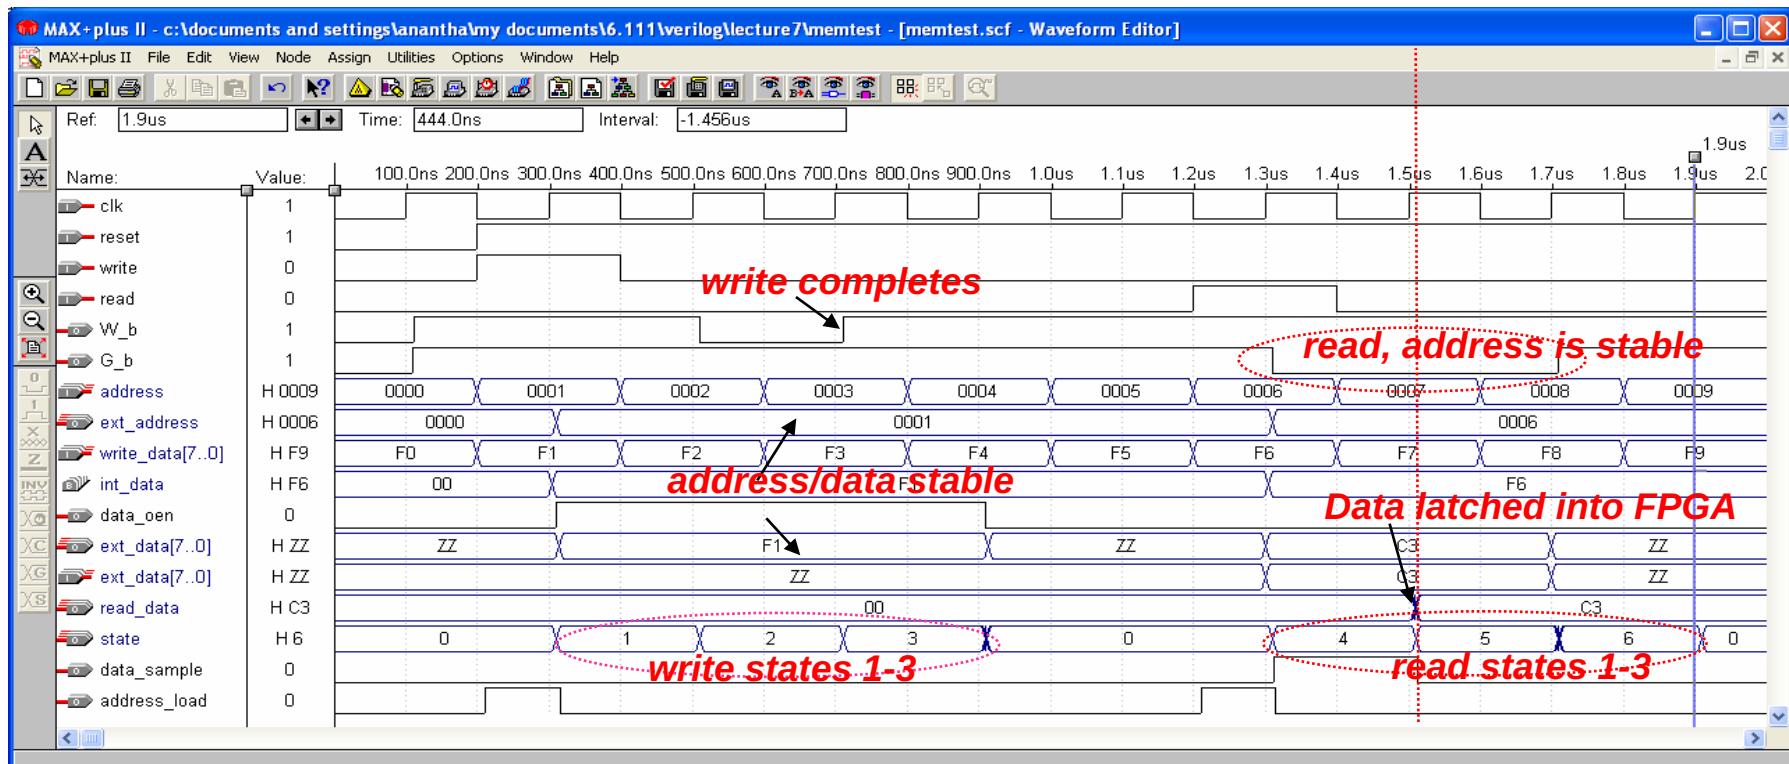
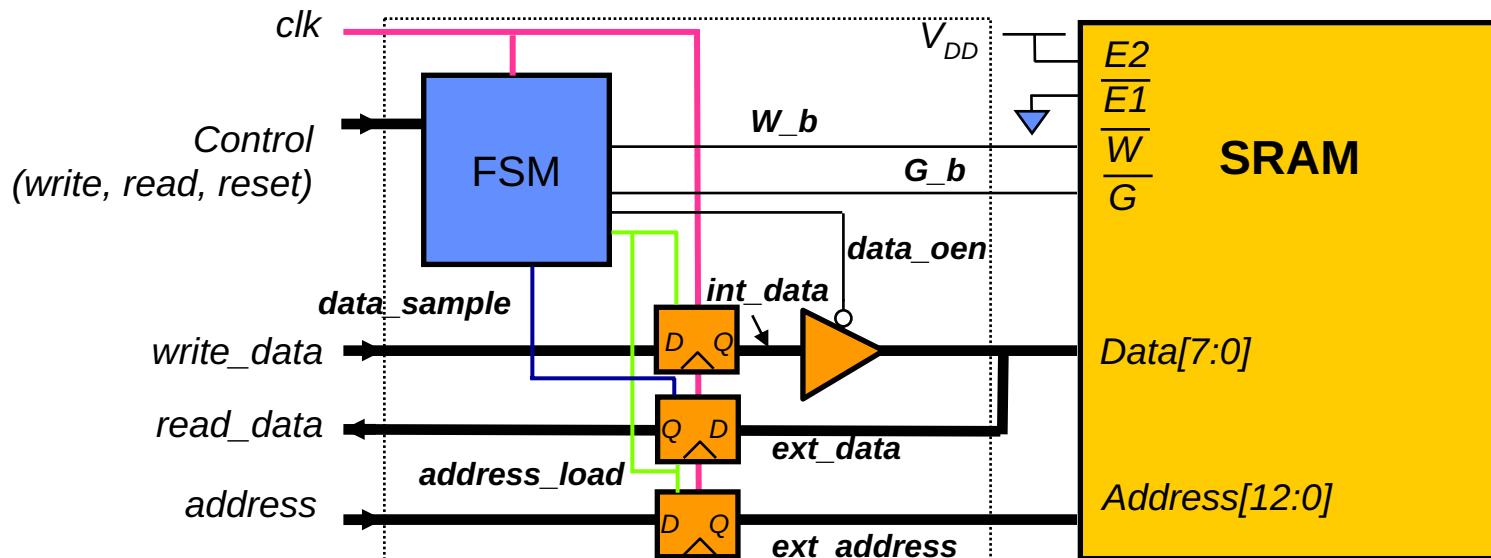
Sample Memory Interface Logic



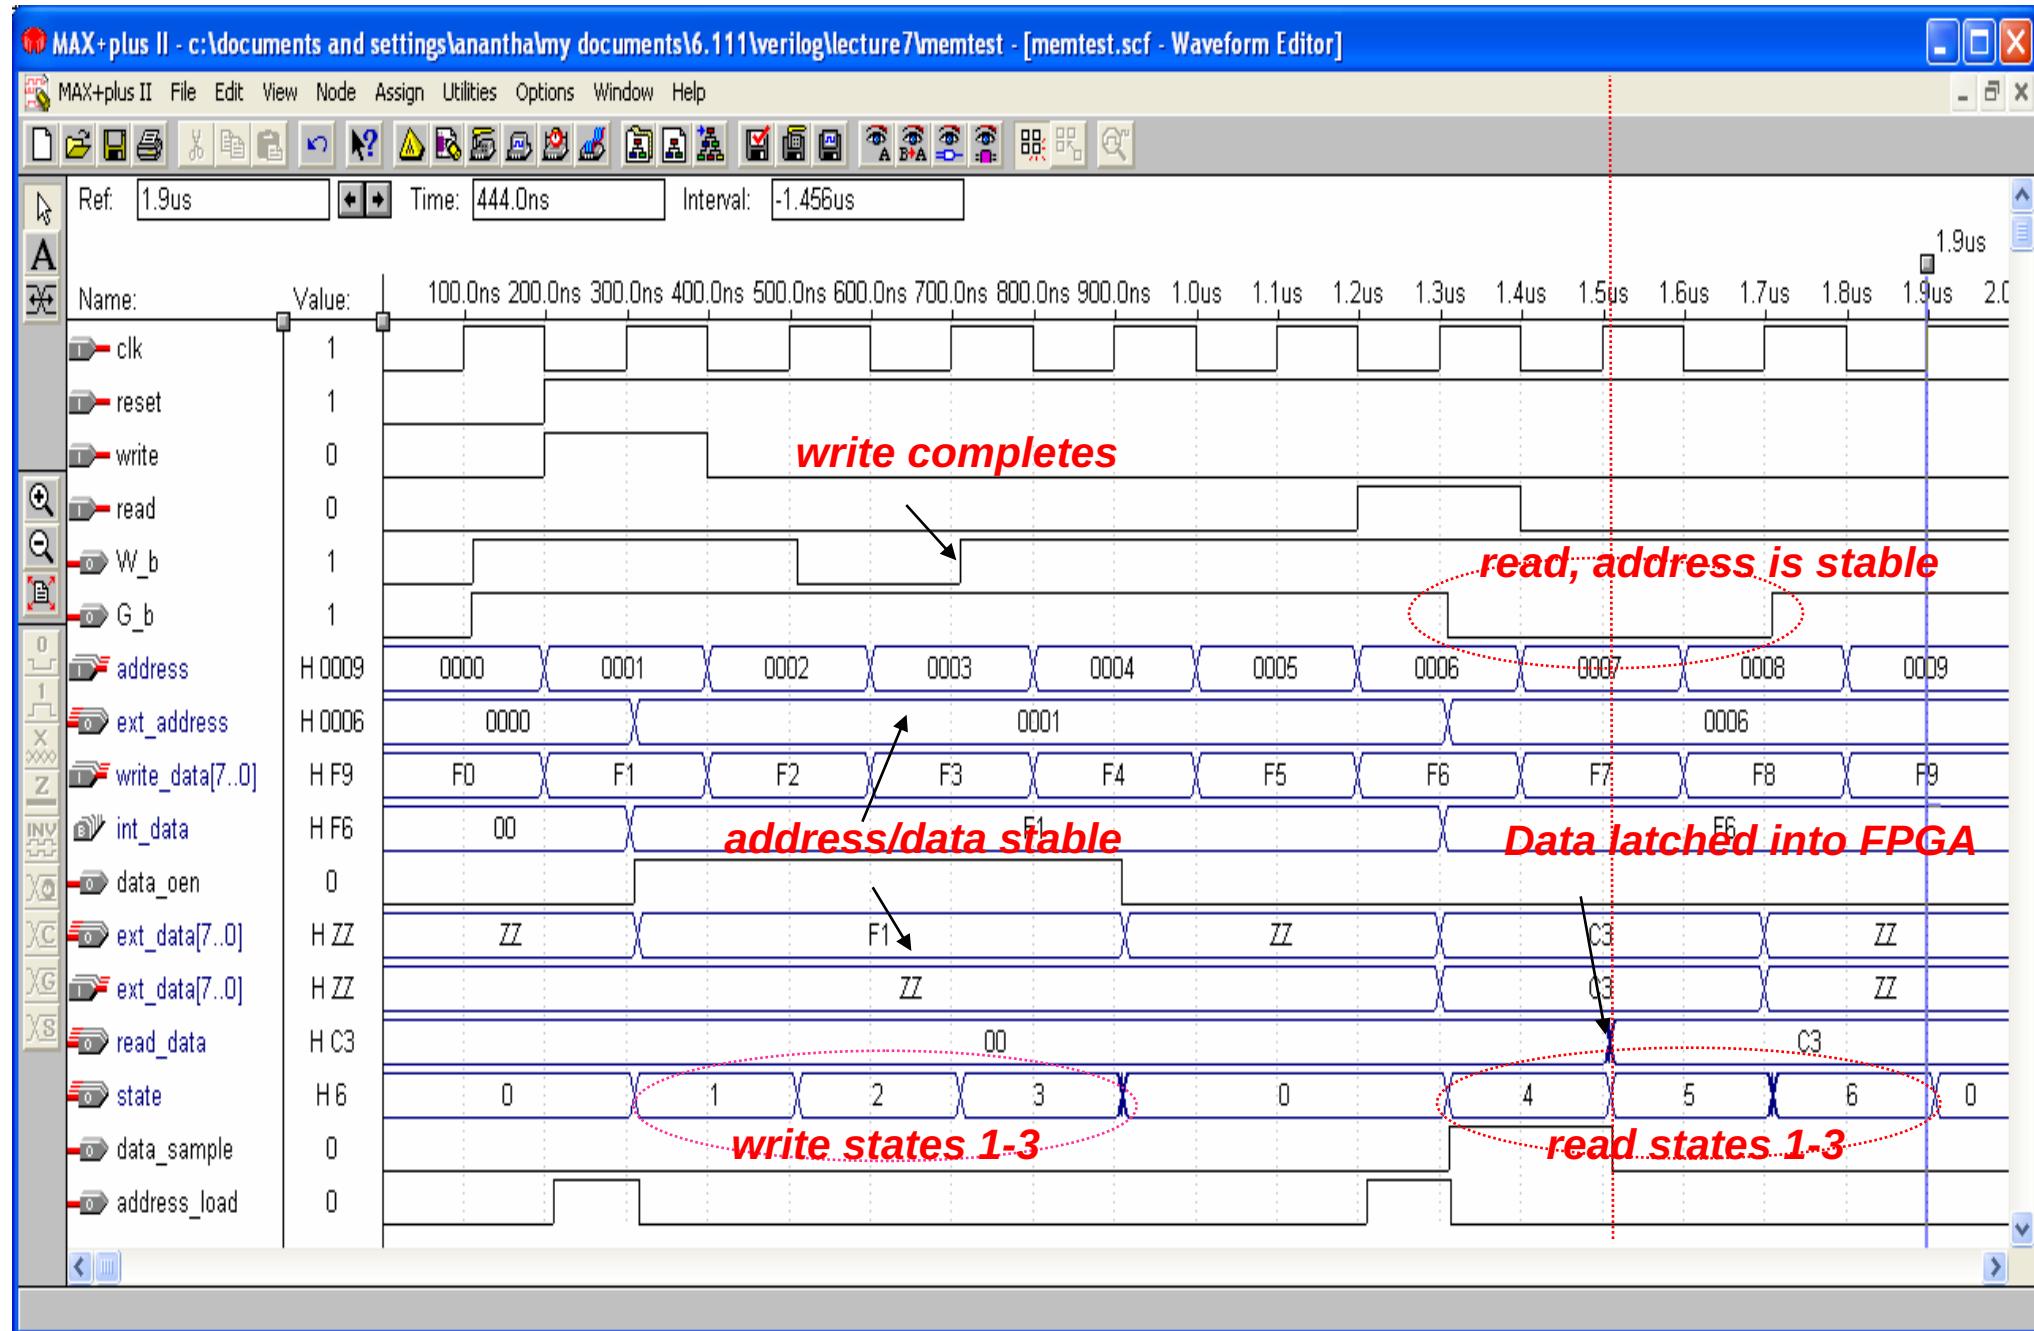
- **Drive data bus only when clock is low**
 - Ensures address and are stable for writes
 - Prevents bus contention
 - Minimum clock period is twice memory access time



Multi-Cycle Read/Write (less aggressive, recommended timing)



Simulation from Previous Slide





Verilog for Simple Multi-Cycle Access



```
module memtest (clk, reset, G_b, W_b, address,
    ext_address, write_data, read_data, ext_data, read,
    write, state, data_oen, address_load, data_sample);
    input clk, reset, read, write;
    output G_b, W_b;
    output [12:0] ext_address;
    reg [12:0] ext_address;
    input [12:0] address;
    input [7:0] write_data;
    output [7:0] read_data;
    reg [7:0] read_data;
    inout [7:0] ext_data;
    reg [7:0] int_data;
    output [2:0] state;
    reg [2:0] state, next;
    output data_oen, address_load, data_sample;
    reg G_b, W_b, G_b_int, W_b_int, address_load,
        data_oen, data_oen_int, data_sample;

    wire [7:0] ext_data;
    parameter IDLE = 0;
    parameter write1 = 1;
    parameter write2 = 2;
    parameter write3 = 3;
    parameter read1 = 4;
    parameter read2 = 5;
    parameter read3 = 6;
```

1/4

// Sequential always block for state assignment

```
assign ext_data = data_oen ? int_data : 8'hz;
```

```
always @ (posedge clk)
begin
    if (!reset) state <= IDLE;
    else state <= next;
```

```
G_b <= G_b_int;
W_b <= W_b_int;
data_oen <= data_oen_int;
if (address_load) ext_address <= address;
if (data_sample) read_data <= ext_data;
if (address_load) int_data <= write_data;
end
```

// note that address_load and data_sample are not
// registered signals

2/4

Verilog for Simple Multi-Cycle Access



```
// Combinational always block for next-state
// computation

always @ (state or read or write) begin
    W_b_int = 1;
    G_b_int = 1;
    address_load = 0;
    data_oen_int = 0;
    data_sample = 0;
    case (state)
        IDLE: if (write) begin
            next = write1;
            address_load = 1;
            data_oen_int = 1;
            end
        else if (read) begin
            next = read1;
            address_load = 1;
            G_b_int = 0;
            end
        else next = IDLE;
    write1: begin
        next = write2;
        W_b_int = 0;
        data_oen_int =1;
        end
    end
```

*Setup the
Default values*

3/4

```
write2: begin
    next = write3;
    data_oen_int =1;
end
write3: begin
    next = IDLE;
    data_oen_int = 0;
end
read1: begin
    next = read2;
    G_b_int = 0;
    data_sample = 1;
end
read2: begin
    next = read3;
    end
read3: begin
    next = IDLE;
    end
default: next = IDLE;
endcase
end
endmodule
```

4/4



Testing Memories



- Common device problems
 - Bad locations: rare for individual locations to be bad
 - Slow (out-of-spec) timing(s): access incorrect data or violates setup/hold
 - Catastrophic device failure: e.g., ESD
 - Missing wire-bonds/devices (!): possible with automated assembly
 - Transient Failures: Alpha particles, power supply glitch
- Common board problems
 - Stuck-at-Faults: a pin shorted to V_{DD} or GND
 - Open Circuit Fault: connections unintentionally left out
 - Open or shorted address wires: causes data to be written to incorrect locations
 - Open or shorted control wires: generally renders memory completely inoperable
- Approach
 - Device problems generally affect the entire chip, almost any test will detect them
 - Writing (and reading back) many different data patterns can detect data bus problems
 - Writing unique data to every location and then reading it back can detect address bus problems

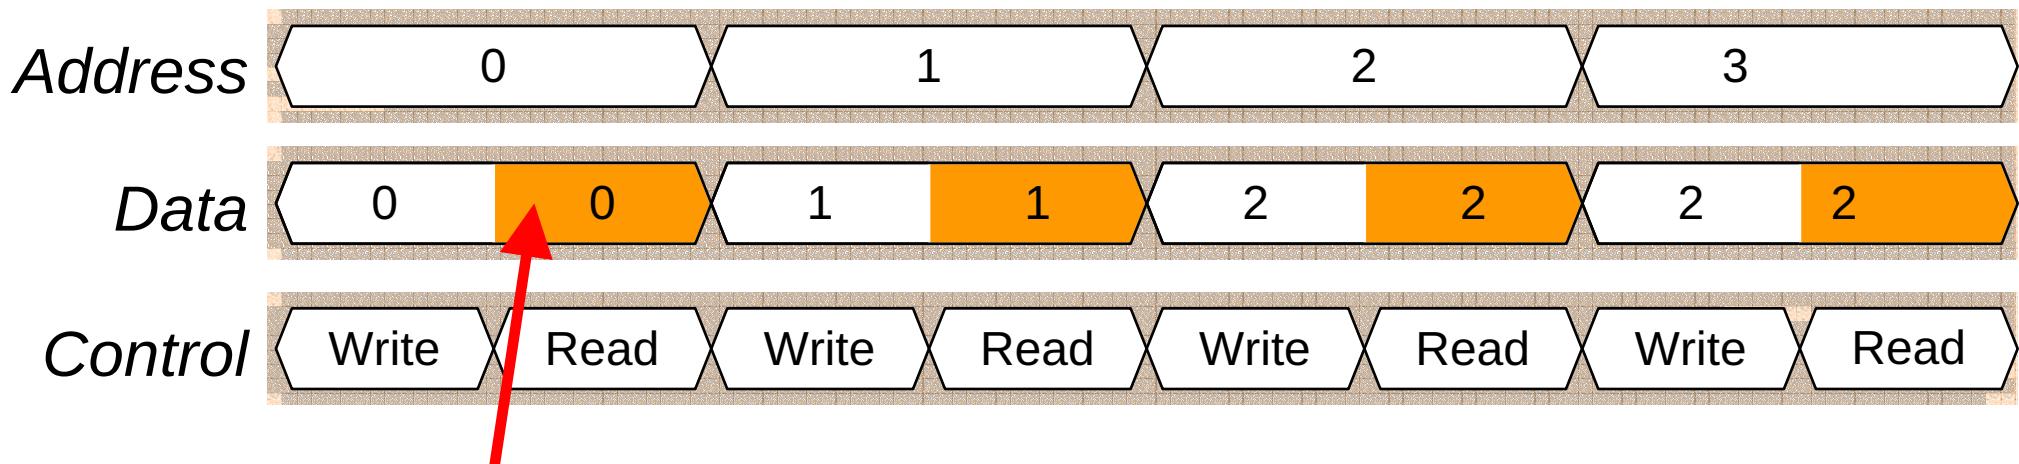
An Approach

■ An idea that almost works

1. Write 0 to location 0
2. Read location 0, compare value read with 0
3. Write 1 to location 1
4. Read location 1, compare value read with 1
5. ...

■ What is the problem?

- Suppose the memory was missing (or output enable was disconnected)



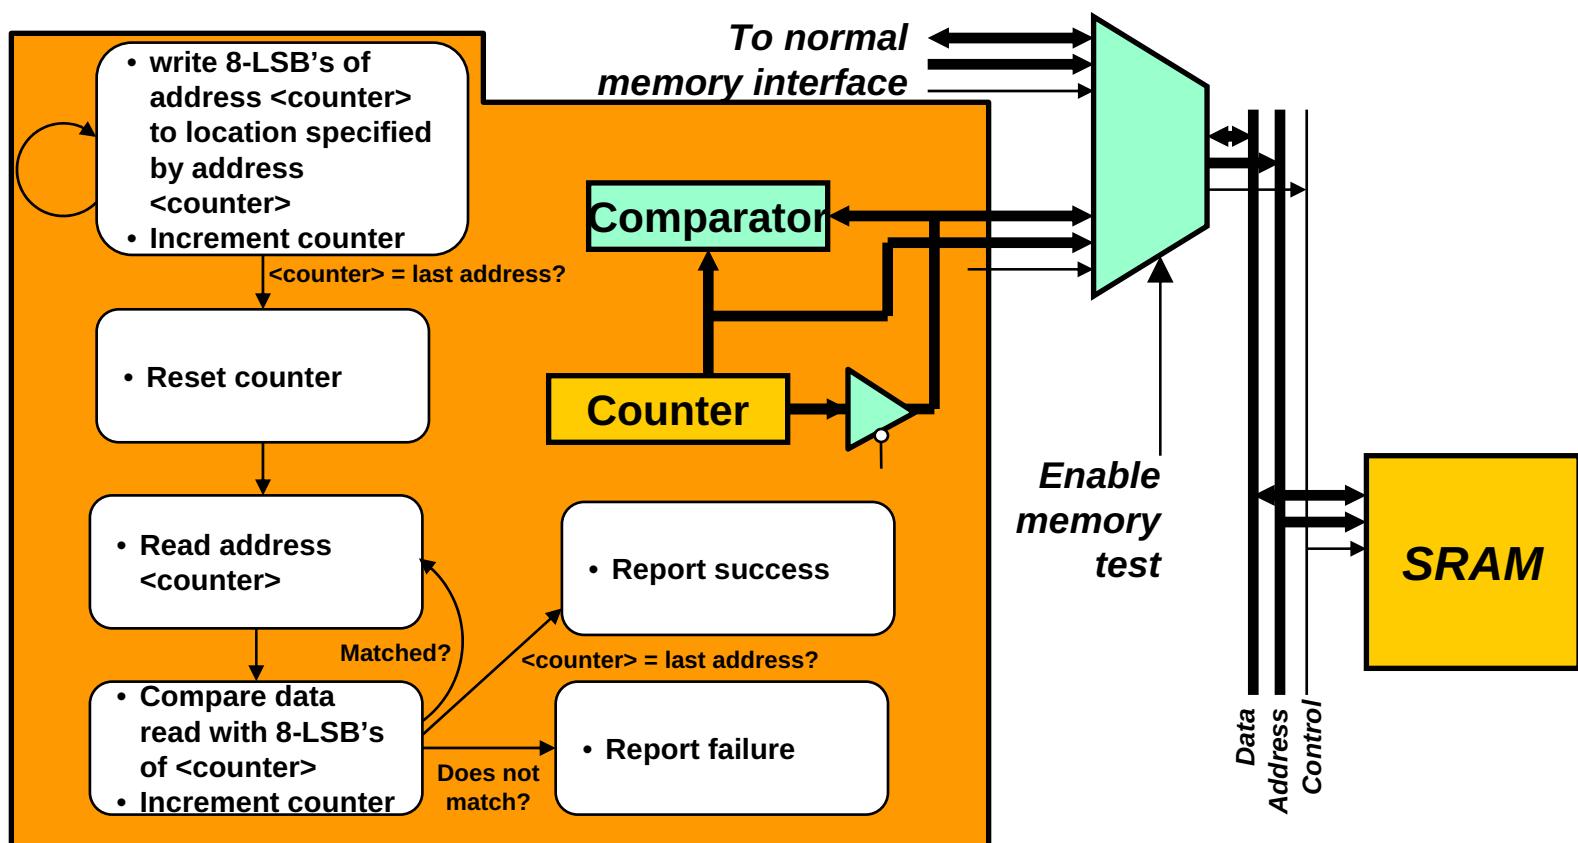
Data bus is undriven but wire capacitance briefly maintains the bus state: memory appears to be ok!

A Simple Memory Tester

- Write to all locations, then read back all locations

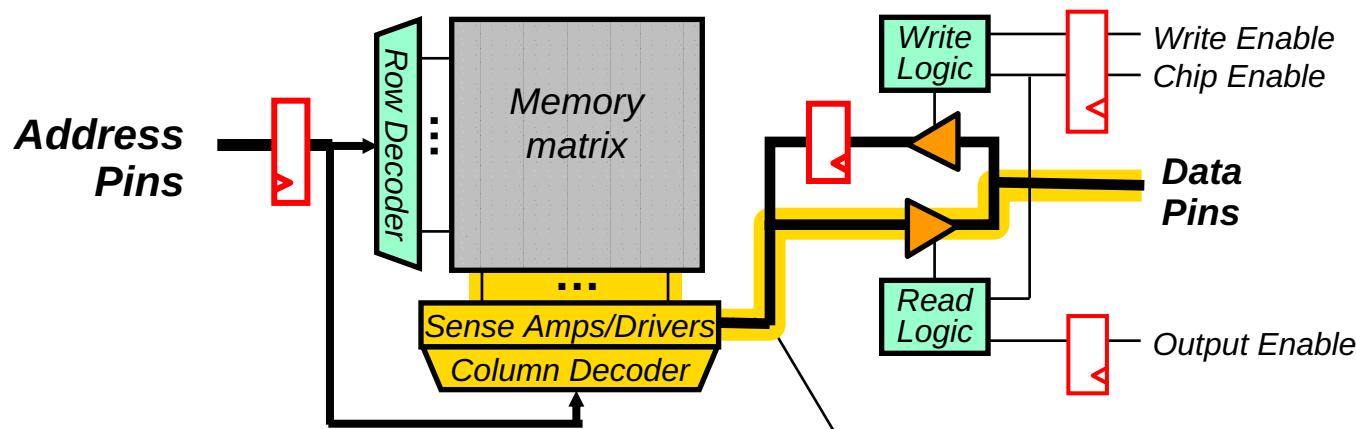
- Separates read/write to the same location with reads/writes of different data to different locations
- (both data and address busses are changed between read and write to same location)

- Write 0 to address 0
- Write 1 to address 1
- ...
- Write $(n \bmod 256)$ to address n
- Read address 0, compare with 0
- Read address 1, compare with 1
- ...
- Read address n, compare with $(n \bmod 256)$



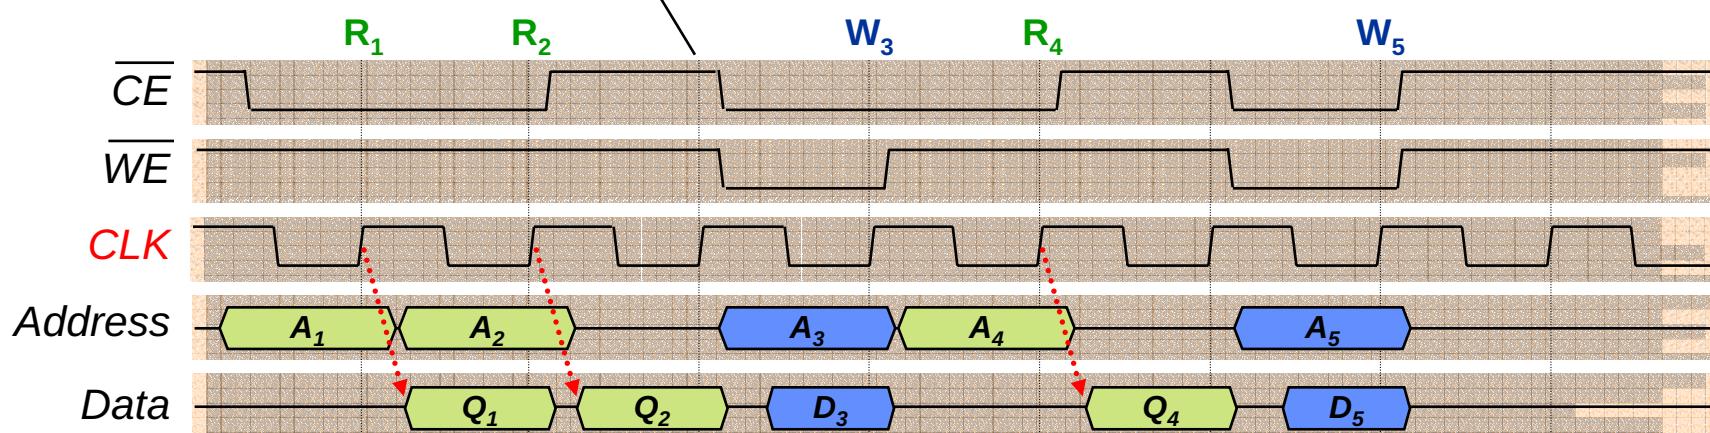
Synchronous SRAM Memories

- Clocking provides input synchronization and encourages more reliable operation at high speeds



difference between read and write timings creates wasted cycles (“wait states”)

long “flow-through” combinational path creates high CLK-Q delay

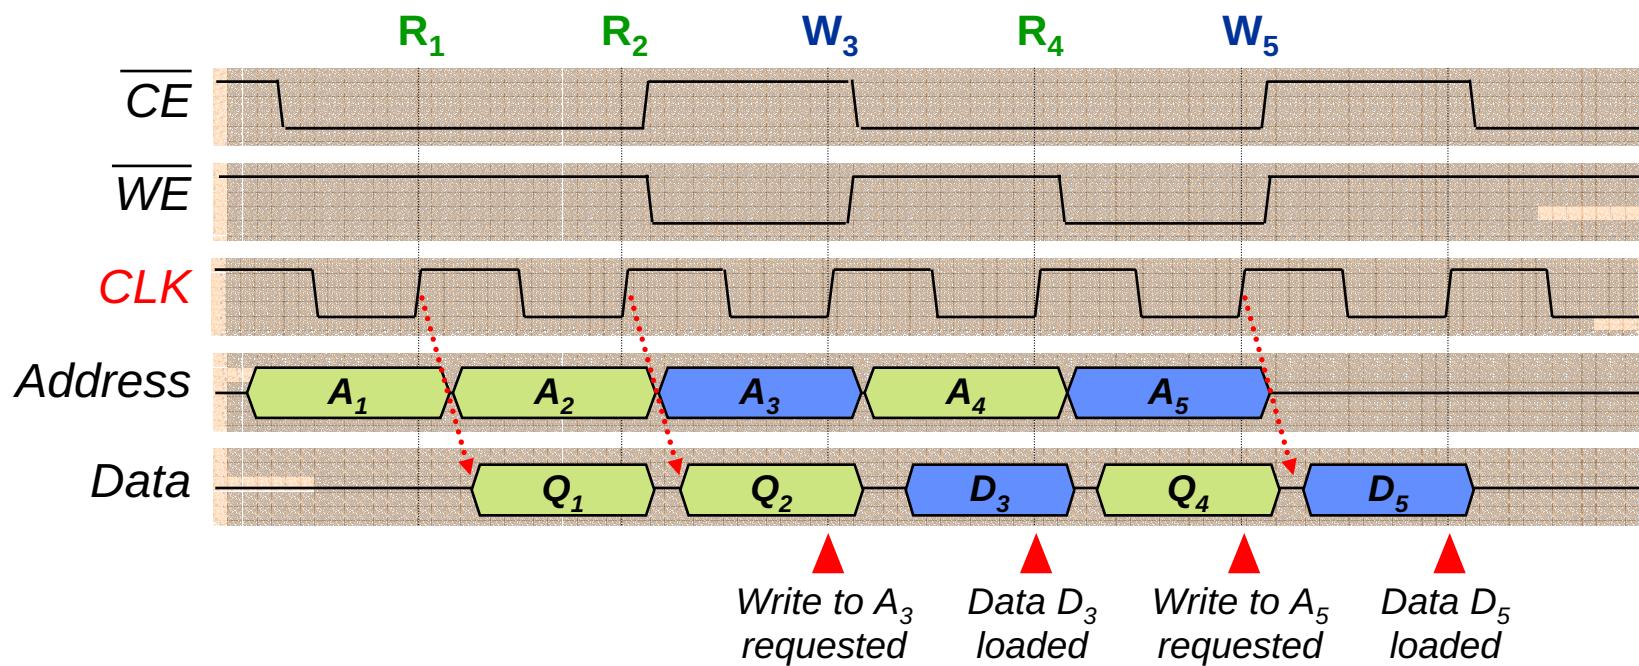




ZBT Eliminates the Wait State

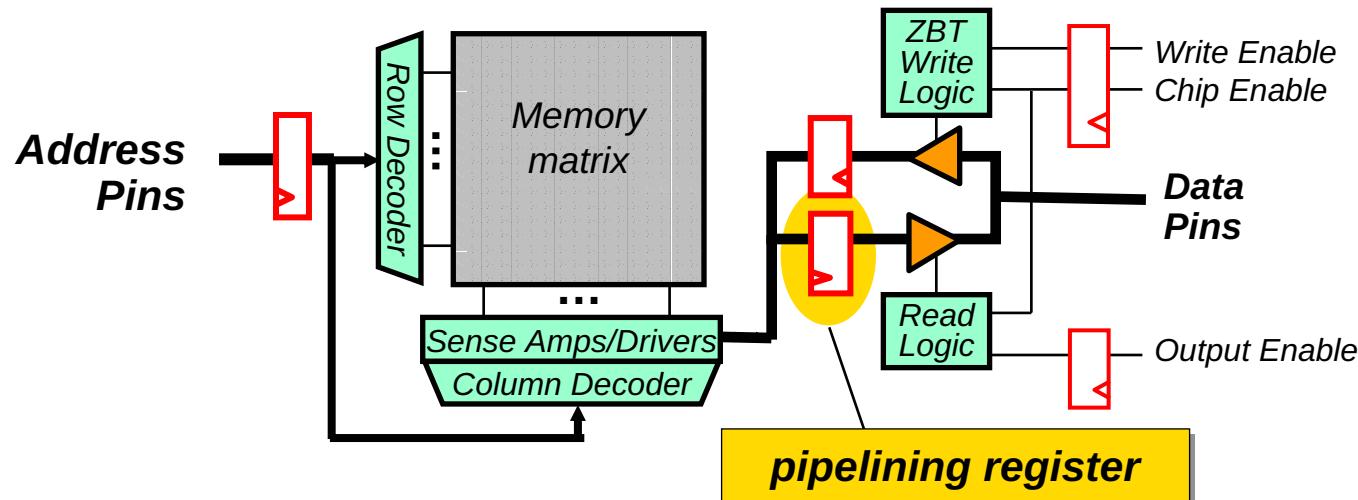


- The wait state occurs because:
 - On a read, data is available *after* the clock edge
 - On a write, data is set up *before* the clock edge
- ZBT (“zero bus turnaround”) memories **change the rules for writes**
 - On a write, data is set up **after** the clock edge
(so that it is read on the following edge)
 - Result: no wait states, higher memory throughput

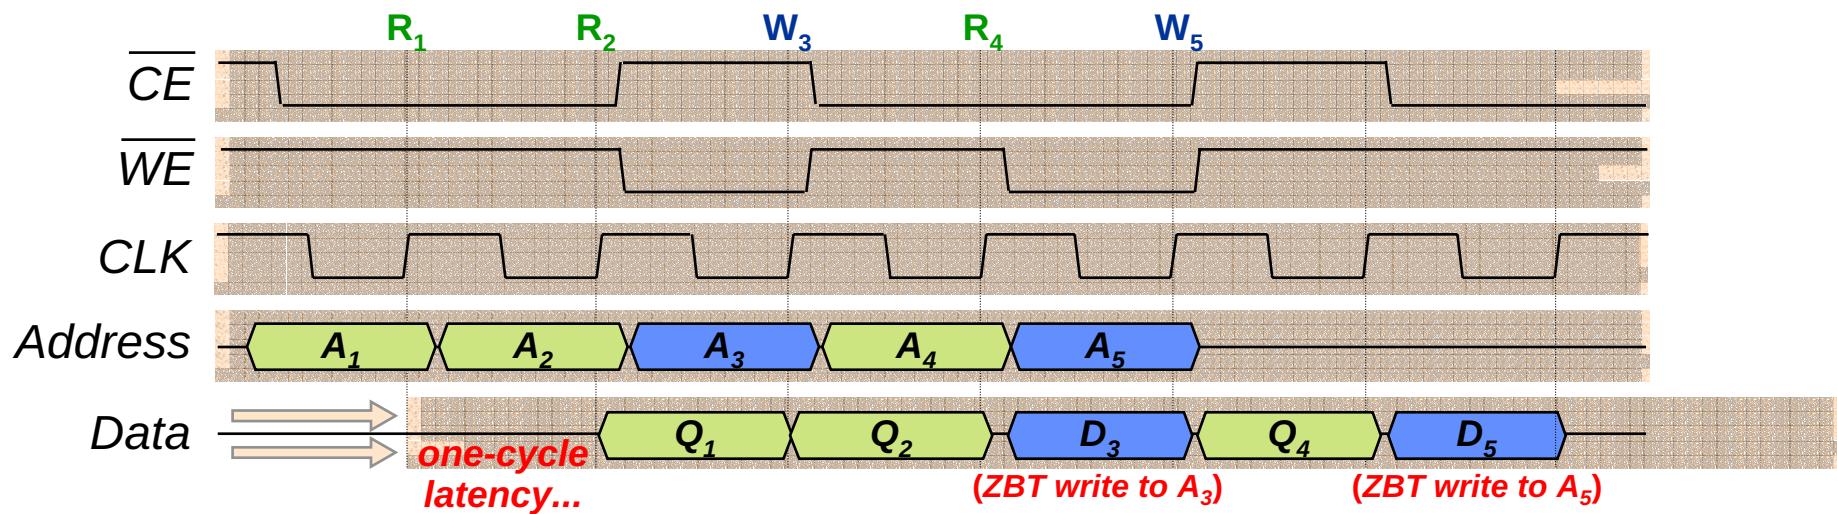


Pipelining Allows Faster CLK

- Pipeline the memory by registering its output
 - Good: Greatly reduces CLK-Q delay, allows higher clock (more throughput)
 - Bad: Introduces an extra cycle before data is available (more latency)

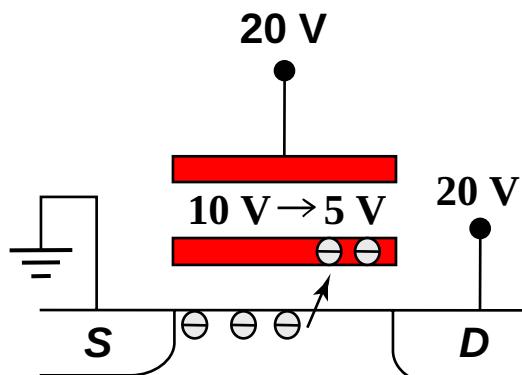


As an example, see
the CY7C147X ZBT
Synchronous SRAM

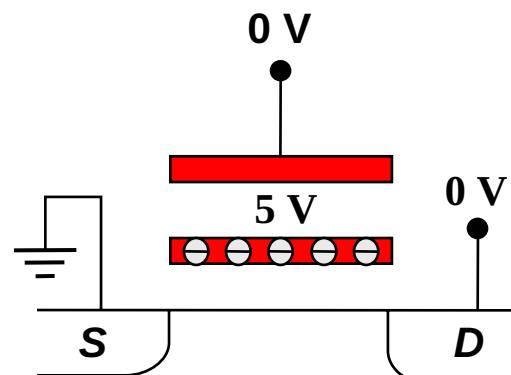




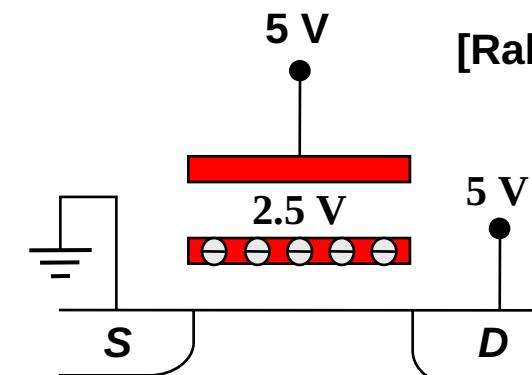
EPROM Cell – The Floating Gate Transistor



Avalanche injection

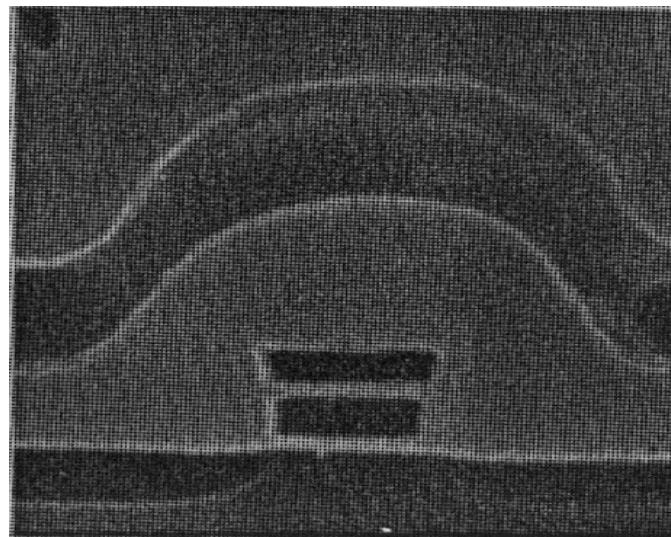


Removing programming voltage leaves charge trapped



Programming results in higher V_T .

[Rabaey03]



EPROM Cell

Courtesy Intel

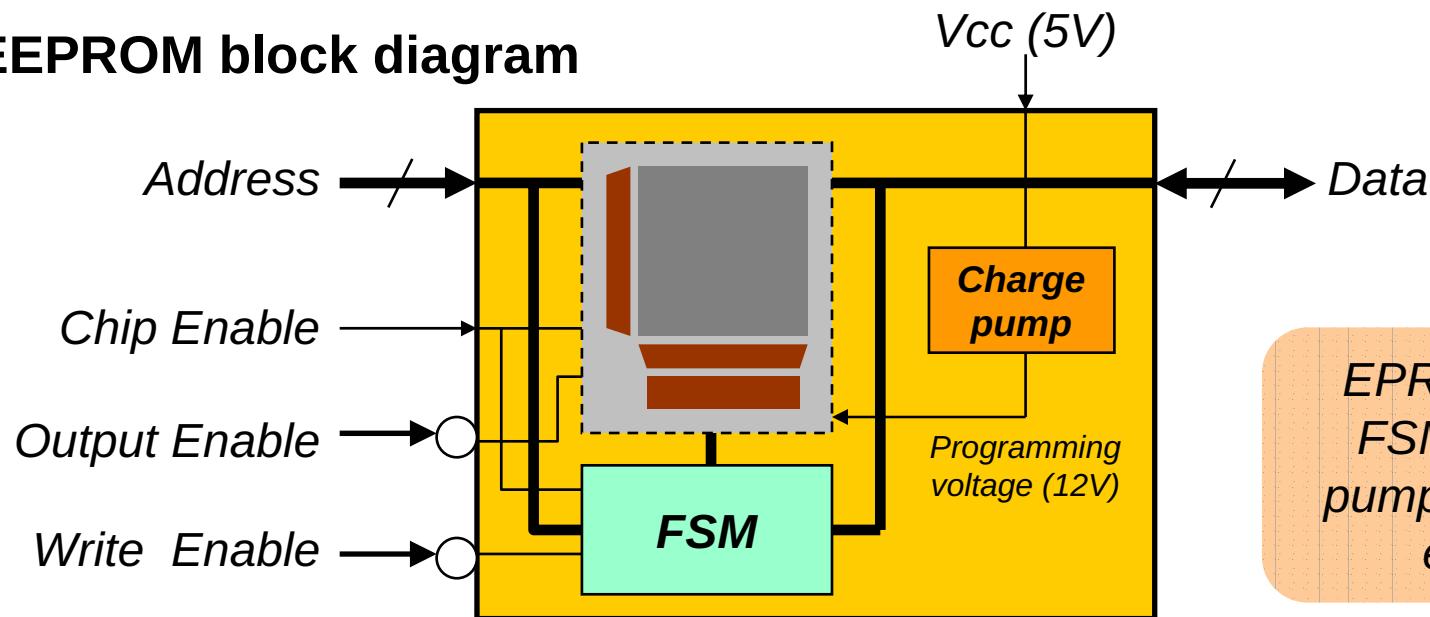
This is a non-volatile memory (retains state when supply turned off)

Interacting with Flash and (E)EPROM



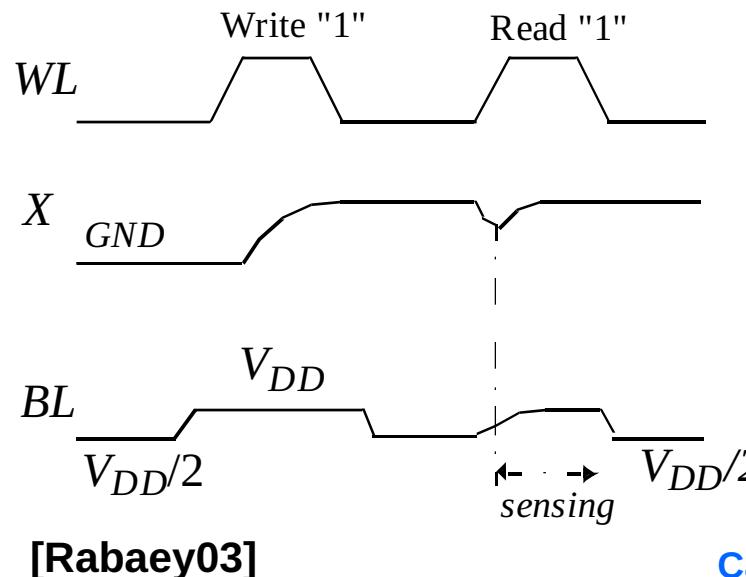
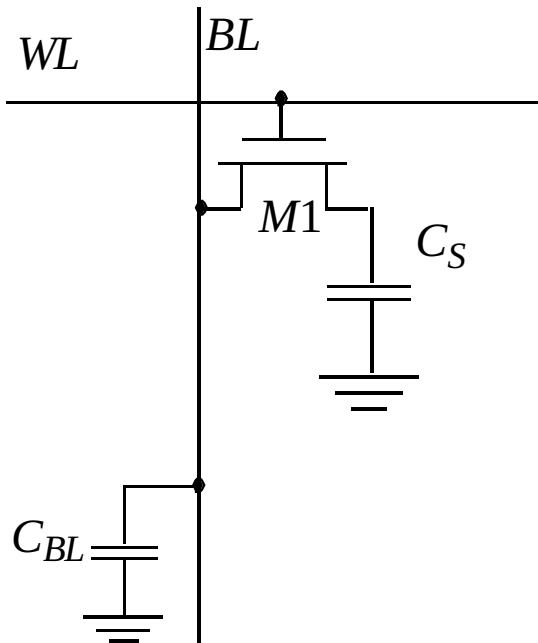
- Reading from flash or (E)EPROM is the same as reading from SRAM
- V_{pp}: input for programming voltage (12V)
 - EPROM: V_{pp} is supplied by programming machine
 - Modern flash/EEPROM devices generate 12V using an on-chip charge pump
- EPROM lacks a write enable
 - Not in-system programmable (must use a special programming machine)
- For flash and EEPROM, write sequence is controlled by an internal FSM
 - Writes to device are used to send signals to the FSM
 - Although the same signals are used, one can't write to flash/EEPROM in the same manner as SRAM

Flash/EEPROM block diagram



EPROM omits
FSM, charge
pump, and write
enable

Dynamic RAM (DRAM) Cell

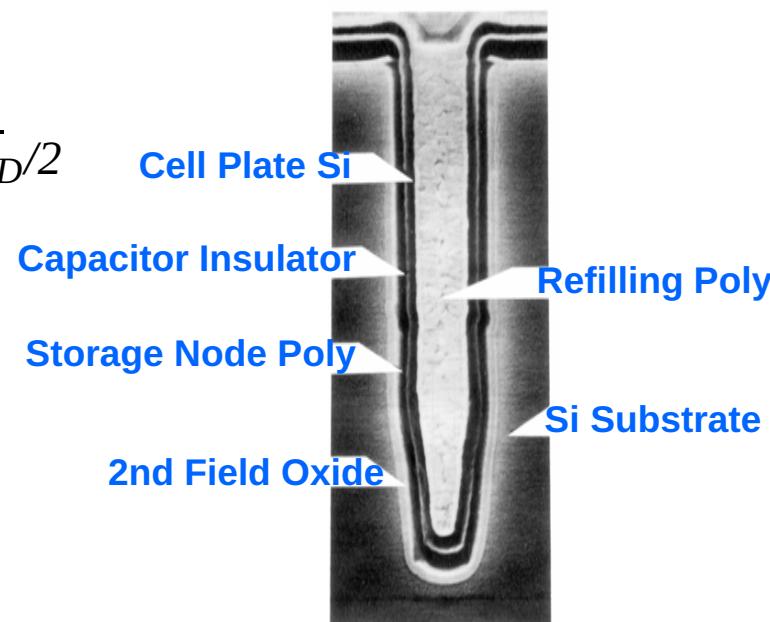


[Rabaey03]

To Write: set Bit Line (BL) to 0 or V_{DD}
& enable Word Line (WL) (i.e., set to V_{DD})

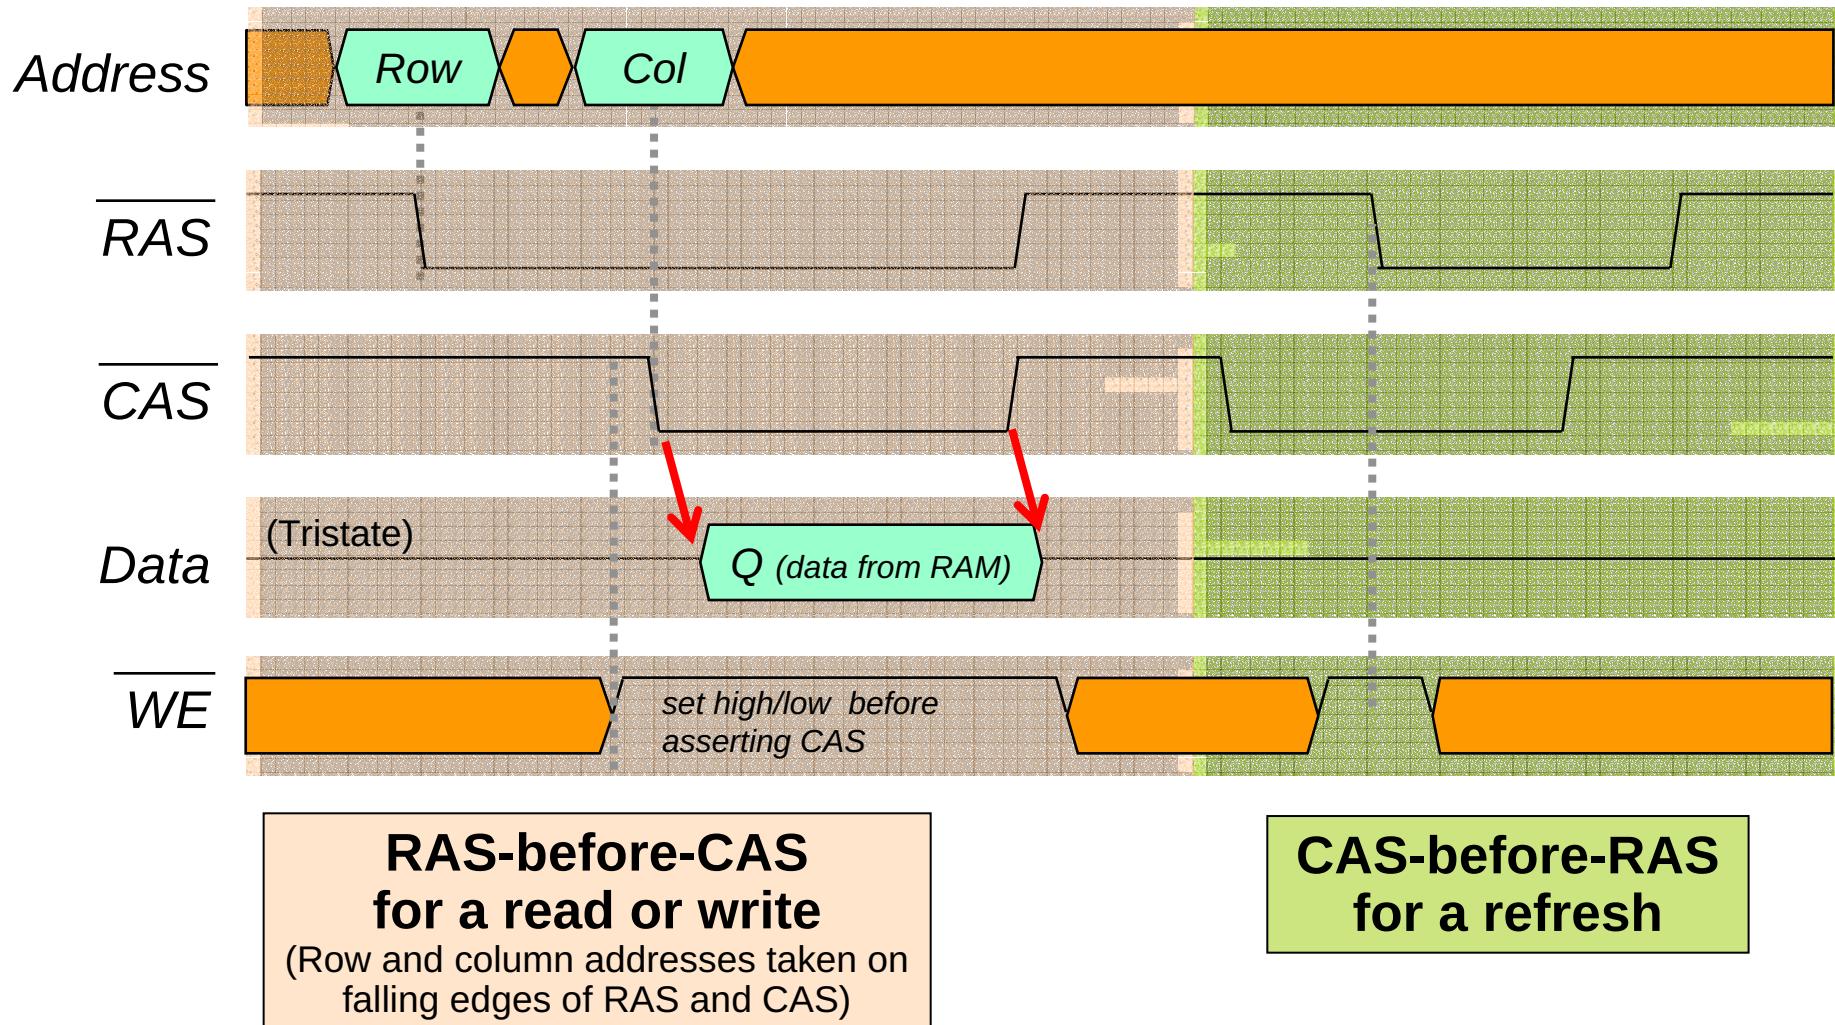
To Read: set Bit Line (BL) to $V_{DD}/2$
& enable Word Line (i.e., set it to V_{DD})

DRAM uses
Special
Capacitor
Structures



- DRAM relies on charge stored in a capacitor to hold state
- Found in all high density memories (one bit/transistor)
- Must be “refreshed” or state will be lost – high overhead

Asynchronous DRAM Operation

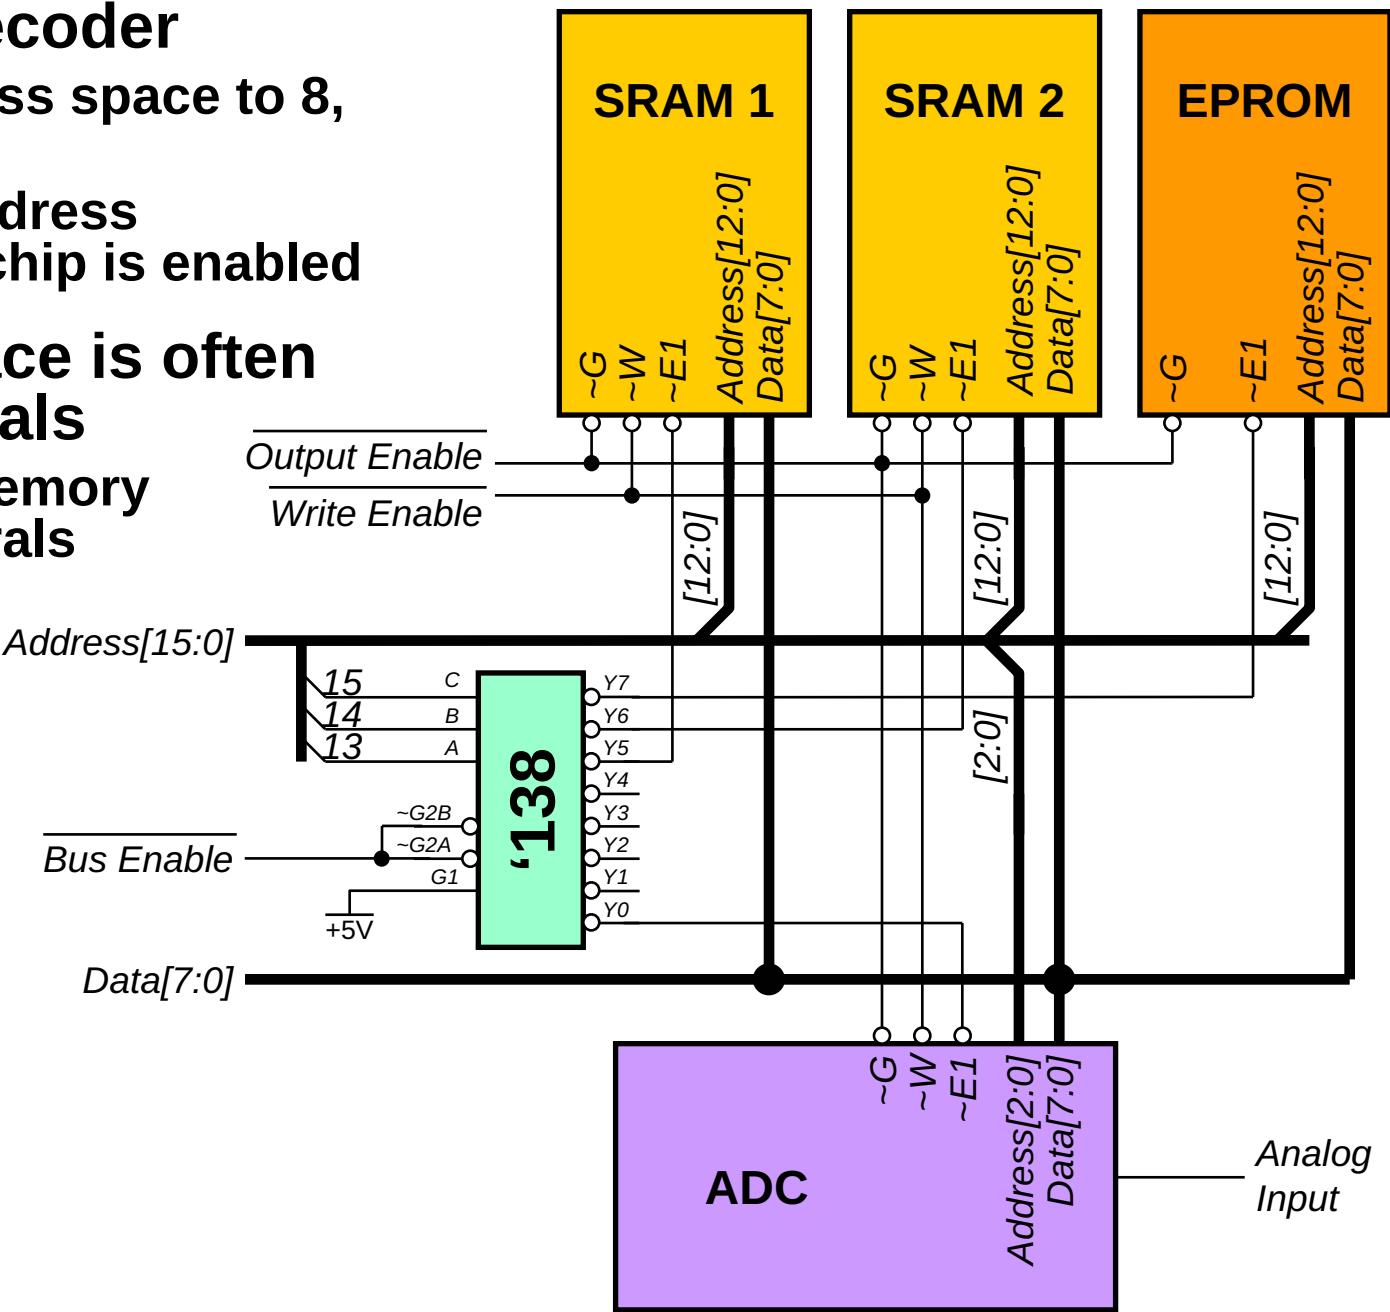
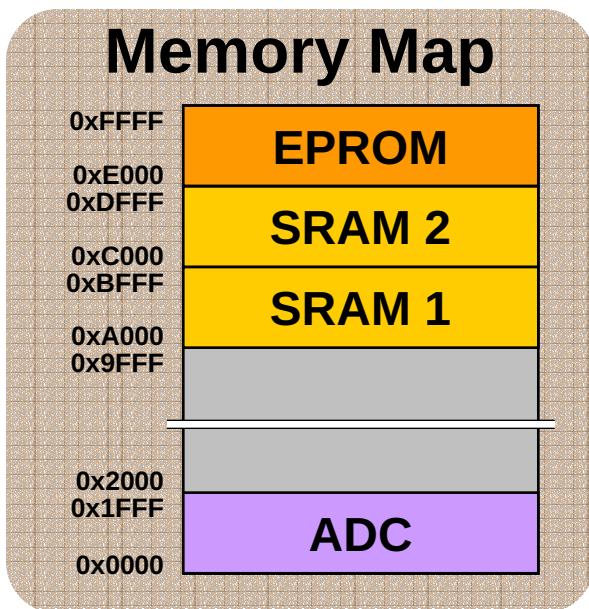


- Clever manipulation of RAS and CAS after reads/writes provide more efficient modes: early-write, read-write, hidden-refresh, etc. (See datasheets for details)

Addressing with Memory Maps



- '138 is a 3-to-8 decoder
 - Maps 16-bit address space to 8, 13-bit segments
 - Upper 3-bits of address determine which chip is enabled
- SRAM-like interface is often used for peripherals
 - Referred to as “memory mapped” peripherals





Key Messages on Memory Devices



■ **SRAM vs. DRAM**

- SRAM holds state as long as power supply is turned on. DRAM must be “refreshed” – results in more complicated control
- DRAM has much higher density, but requires special capacitor technology.
- FPGA usually implemented in a standard digital process technology and uses SRAM technology

■ **Non-Volatile Memory**

- Fast Read, but very slow write (EPROM must be removed from the system for programming!)
- Holds state even if the power supply is turned off

■ **Memory Internals**

- Has quite a bit of analog circuits internally -- pay particular attention to noise and PCB board integration

■ **Device details**

- Don't worry about them, wait until 6.012 or 6.374